

PAC Learning

Machine Learning 2023-24

UML Book Chapter 3

Slides P. Zanuttigh (some material from F. Vandin slides)


Supervised Learning




Training data
with labels



Training procedure



Data to be
analyzed



ML model

(training: estimate parameters)



Predicted
Label

In most of the course we will focus on supervised learning

Probably Approximately Correct (PAC) learning

Since the training data is sampled accordingly to D :

- ❑ we can only be approximately correct
- ❑ we can only be probably correct

Parameters:

- ❑ accuracy parameter ϵ : we are satisfied with a good h_S for which $L_{D,f}(h_S) \leq \epsilon$
- ❑ confidence parameter δ : want h_S to be a good hypothesis with probability $\geq 1 - \delta$

Theorem


(Finite Hypothesis Classes are PAC Learnable)

Let H be a finite hypothesis class. Let $\delta \in (0,1)$, $\epsilon \in (0,1)$ and $m \in \mathbb{N}$ such that:

$$m \geq \frac{\log\left(\frac{|H|}{\delta}\right)}{\epsilon}$$

Notice: m grows with $|H|$ and is inversely proportional to δ and ϵ

Then for any f and any D for which the realizability assumption holds, with probability $\geq 1 - \delta$ we have that for every ERM hypothesis h_s it holds that



m : size of the training set (i.e., S contains m i.i.d. samples)

Definition: PAC Learnability

A hypothesis class \mathcal{H} is **PAC learnable** if there exist a function $m_{\mathcal{H}}: (0,1)^2 \rightarrow \mathbb{N}$ and a learning algorithm such that for **every** $\delta, \epsilon \in (0,1)$, for **every** distribution D over \mathcal{X} and for **every** labeling function $f: \mathcal{X} \rightarrow \{0,1\}$, **if the realizability assumption holds** with respect to \mathcal{H}, D, f when running the algorithm on $m \geq m_{\mathcal{H}}(\epsilon, \delta)$ i.i.d. examples generated by D and labeled by f the algorithm returns a hypothesis h such that, with probability $\geq 1 - \delta$ (over the choice of the m training examples): $L_{D,f}(h) \leq \epsilon$

Notes:

- It is a property of the hypothesis class
- Must be satisfied for **every** δ, ϵ, D, f
- $m_{\mathcal{H}}: (0,1)^2 \rightarrow \mathbb{N}$: sample complexity of learning \mathcal{H} (depends on ϵ and δ)
- $m_{\mathcal{H}}$ is the minimum integer that satisfies the requirements
- **Sample complexity**: minimum size of training set to be sure to satisfy requirements (**sufficient but not necessary condition**)
- Only if **realizability assumption** holds
- PAC: Good probability ($1 - \delta$) of having a good predictor ($L_{D,f}(h) \leq \epsilon$)

Corollary (PAC)

A hypothesis class \mathcal{H} is **PAC learnable** if there exist a function $m_{\mathcal{H}}: (0,1)^2 \rightarrow \mathbb{N}$ and a learning algorithm such that for **every** $\delta, \epsilon \in (0,1)$, for **every** distribution D over \mathcal{X} and for **every** labeling function $f: \mathcal{X} \rightarrow \{0,1\}$, if the **realizability assumption** holds with respect to \mathcal{H}, D, f when running the algorithm on $m \geq m_{\mathcal{H}}(\epsilon, \delta)$ i.i.d. examples generated by D and labeled by f the algorithm returns a hypothesis h such that, with probability $\geq 1 - \delta$ (over the choice of the m training examples): $L_{D,f}(h) \leq \epsilon$

Let H be a finite hypothesis class. Let $\delta \in (0,1)$, $\epsilon \in (0,1)$ and $m \in \mathbb{N}$ such that:

$$m \geq \frac{\log\left(\frac{|H|}{\delta}\right)}{\epsilon}$$

Then for any f and any D for which the realizability assumption holds, with probability $\geq 1 - \delta$ we have that for every ERM hypothesis h_s it holds that

$$L_{D,f}(h_s) \leq \epsilon$$

probably approximately correct

From the previous theorem and from the definition of PAC:

Corollary (PAC)

Every **finite** hypothesis class is PAC learnable with sample complexity

$$m_H(\epsilon, \delta) \leq \left\lceil \frac{\log\left(\frac{|H|}{\delta}\right)}{\epsilon} \right\rceil$$

- Larger hypothesis class $|H| \rightarrow$ need a larger training set
- Smaller training error $\epsilon \rightarrow$ need a larger training set
- Better probability of a good solution $1 - \delta \rightarrow$ need a larger training set

sufficient condition
(not necessary)

Drop Some Assumptions

1. **Realizability Assumption**: there exists $h^* \in \mathcal{H}$ such that

$$L_{D,f}(h) = 0$$
 - ❑ Too strong in many real-world applications! *no “perfect” predictor*
2. **Function f** : in many applications it is not too realistic that the *labeling is fully determined* by the features we measure
 - ❑ *Relaxation*: replace target labeling function...
 - ❑ .. and assume D is a *probability distribution over $\mathcal{X} \times \mathcal{Y}$*
 - i.e., D is the joint distribution over domain points and labels
 - ❑ For example, two components of D :
 - D_x : (marginal) distribution over domain points
 - $D((x, y)|x)$: conditional distribution over labels for each domain point
3. **Binary classification** → move to a more general setting
 - ❑ Multi-class classification and regression problems



$(w, h) \rightarrow M \text{ or } F ??$

Empirical and True Error

Assuming D is a probability distribution over $\mathcal{X} \times \mathcal{Y}$, the *true error* (or risk) is:

$$L_D(h) \triangleq \mathbb{P}_{(x,y) \sim D} [h(x) \neq y] \stackrel{\text{def}}{=} D(\{(x, y): h(x) \neq y\})$$



Recall: previously $\mathbb{P}_{x \sim D} [h(x) \neq f(x)]$

As before D is not known to the ML algorithm (learner): the learner only knows the training data S

The *Empirical Risk* is as before:

$$L_S(h) \stackrel{\text{def}}{=} \frac{|\{i: h(x_i) \neq y_i, 1 \leq i \leq m\}|}{m}$$

Note: $L_S(h)$ = probability for a pair taken uniformly (x_i, y_i) at random from S the event $h(x_i) \neq y_i$ holds


Bayes Optimal Predictor

Learner's goal: find $h: \mathcal{X} \rightarrow \mathcal{Y}$ minimizing $L_D(h)$

Question: Is there a best predictor (for binary classification) ?

Given a probability distribution D over $\mathcal{X} \times \{0,1\}$, the best predictor is the *Bayes Optimal Predictor*:

$$f_D(x) = \begin{cases} 1 & \text{if } P[y = 1 | x] \geq \frac{1}{2} \\ 0 & \text{otherwise} \end{cases}$$

binary classification


Proposition:

For any classifier $g: \mathcal{X} \rightarrow \{0,1\}$, it holds $L_D(f_D) \leq L_D(g)$

Can we use
such predictor ?

D (and consequently $\mathbb{P}[y = 1 | x]$) is not known!

Definition: Agnostic PAC Learnability

- ❑ *Idea*: We drop the requirement of finding the best predictor, but we do not want to be too far from it
- ❑ Definition (*agnostic* PAC learnability):

A hypothesis class \mathcal{H} is *agnostic* PAC learnable if there exist a function $m_{\mathcal{H}}: (0,1)^2 \rightarrow \mathbb{N}$ and a learning algorithm such that for every $\delta, \epsilon \in (0,1)$ and for every distribution D over $\mathcal{X} \times \mathcal{Y}$, when running the algorithm on $m \geq m_{\mathcal{H}}(\epsilon, \delta)$ i.i.d. examples generated by D the algorithm returns a hypothesis h such that, with probability $\geq 1 - \delta$ (over the choice of the m training examples):

$$L_D(h) \leq \min_{h' \in \mathcal{H}} L_D(h') + \epsilon$$

Not farther than ϵ from
the best predictor in \mathcal{H}

- ❑ This is a generalization of the previous learning model
- ❑ *Realizability* is not required

Multiclass Classification and Regression

We consider 3 possible learning problems:

- Domain set and training data have the same structure for all problems
- Learner output: $h: \mathcal{X} \rightarrow \mathcal{Y}$ (\mathcal{Y} is different for the 3 problems)

1. **Binary classification**: (it is the one we were considering before)

- Target set: $\mathcal{Y} = \{0, 1\}$ (target set size = 2)



2. **Multiclass classification** with $K > 2$ classes

- Target set: $\mathcal{Y} = \{0, 1, \dots, K - 1\}$ (target set size = K)
- Loss: as for the binary case



3. **Regression** $\mathcal{Y} = \mathbb{R}$

- Target set has an infinite size
- Need a new loss function !



10M€



300k€



50k€

Generalized Loss Function

Given:

- \mathcal{H} : hypothesis class
- Z : domain $(\mathcal{X} \times \mathcal{Y})$

A loss function is any function $l: \mathcal{H} \times Z \rightarrow \mathbb{R}_+$

Risk function: expected loss of an hypothesis $h \in \mathcal{H}$ with respect to D over Z :

$$L_D(h) \stackrel{\text{def}}{=} \mathbb{E}_{Z \sim D}[l(h, z)]$$

Empirical risk: expected loss over a given training set

$S = (z_1, \dots, z_m) \in Z^m$:

$$L_S(h) \stackrel{\text{def}}{=} \frac{1}{m} \sum_{i=1}^m l(h, z_i)$$

Common Loss Functions

- ❑ *0-1 loss*: Commonly used in binary or multiclass **classification**

$$l_{0-1}(h, (x, y)) \stackrel{\text{def}}{=} \begin{cases} 0 & \text{if } h(x) = y \\ 1 & \text{if } h(x) \neq y \end{cases}$$

- ❑ *Cross Entropy* is also used in classification (presented later in the course)
- ❑ *Squared loss* (**L2**): Commonly used in **regression**, penalize few large errors

$$l_{sq}(h, (x, y)) \stackrel{\text{def}}{=} (h(x) - y)^2$$

- ❑ *Absolute value loss* (**L1**): Commonly used in **regression**, penalize many small errors

$$l_{abs}(h, (x, y)) \stackrel{\text{def}}{=} |h(x) - y|$$

- ❑ In general, the loss function depends on the application!
- ❑ But computational considerations must also be taken into account..

Optimal Loss Depends on Application


The relevance of different error types depends on the application

Example: fingerprints classification/verification



Two types of error:

- **False accept**: accept an unauthorized user
- **False reject**: do not accept an authorized user



		f	
		$+1$	-1
h	$+1$	no error	false accept
	-1	false reject	no error

Optimal Loss: Discount Verification


The error measure - for supermarkets

Supermarket verifies fingerprint for discounts

False reject is costly; customer gets annoyed!

False accept is minor; gave away a discount and intruder left their fingerprint 😊

		f	
		+1	-1
h	+1	0	1
	-1	10	0



$\begin{cases} +1 & \text{you} \\ -1 & \text{intruder} \end{cases}$



Optimal Loss: CIA Data Center


The error measure - for the CIA


CIA verifies fingerprint for security

False accept is a disaster!

False reject can be tolerated

Try again; you are an employee 😊

		f	
		+1	-1
h	+1	0	1000 
	-1	1	0



$\left\{ \begin{array}{ll} +1 & \text{you} \\ -1 & \text{intruder} \end{array} \right.$



Agnostic PAC Learnability: General Loss Functions

Definition

A hypothesis class \mathcal{H} is *agnostic* PAC learnable with respect to a set Z and a loss function $l: H \times Z \rightarrow \mathbb{R}_+$ if there exist a function $m_{\mathcal{H}}: (0,1)^2 \rightarrow \mathbb{N}$ and a learning algorithm such that for every $\delta, \epsilon \in (0,1)$ and for every distribution D over Z , when running the algorithm on $m \geq m_{\mathcal{H}}(\epsilon, \delta)$ i.i.d. examples generated by D the algorithm returns a hypothesis h such that, with probability $\geq 1 - \delta$ (over the choice of the m training examples):

$$L_D(h) \leq \min_{h' \in \mathcal{H}} L_D(h') + \epsilon$$

where $L_D(h) = \mathbb{E}_{z \sim D}[l(h, z)]$