

# Stochastic Gradient Descent

Machine Learning 2023-24

UML book chapter 14 (the slides contain a simplified presentation)

Slides P. Zanuttigh (some material F. Vandin, S. Rudes)

# Minimize a Differentiable Function

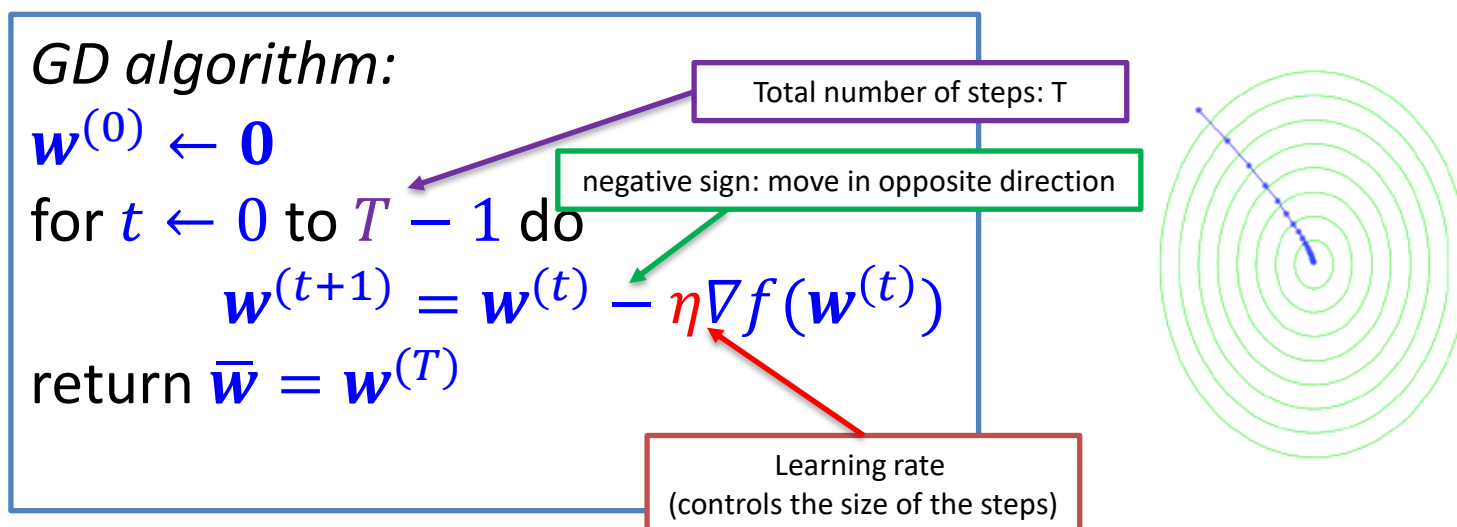
- *The task*: Need a general approach to **minimize** a **differentiable** convex function  $f(\mathbf{w})$  with respect to the (**weights'**) vector  $\mathbf{w}$ 
  - Notice: the dimensionality of  $\mathbf{w}$  can be **very** large
- *Recall*: the gradient  $\nabla f(\mathbf{w})$  of a differentiable function  $f: \mathbb{R}^d \rightarrow \mathbb{R}$  is:

$$\nabla f(\mathbf{w}) = \left( \frac{\partial f(\mathbf{w})}{\partial w_1}, \dots, \frac{\partial f(\mathbf{w})}{\partial w_d} \right)$$

- *Idea*: the gradient points in the direction of the largest increase of  $f$  in the region close to  $\mathbf{w}$
- Move in the opposite direction until you find a minima
- Gradient corresponds to first order Taylor approximation
  - First order Taylor:  $f(\mathbf{u}) \approx f(\mathbf{w}) + \langle \mathbf{u} - \mathbf{w}, \nabla f(\mathbf{w}) \rangle$
  - Good approximation for small steps  $\rightarrow$  need to move **step by step**
- The theory can be extended to non-differentiable functions using subgradients (*if interested see the book, not part of the course*)

# Gradient Descent (GD)

General approach to minimize a differentiable convex function  $f(\mathbf{w})$



- Start from an initial point
  - e.g.,  $\mathbf{w}^{(0)} = \mathbf{0}$  or random value or initial guess (and many other strategies)
- At each step move in direction opposite to the gradient
- Stop when solution does not improve or max iterations reached
- Get the final point or the one corresponding to minimum value of the objective function

# Gradient Descent: Accuracy and Convergence

## *Hypothesis:*

- $f(\mathbf{w})$  is a **convex**  $\rho$ -Lipschitz function
  - recall  $\rho$ -Lipschitz:  $\|f(\mathbf{w}_1) - f(\mathbf{w}_2)\| \leq \rho \|\mathbf{w}_1 - \mathbf{w}_2\|$
- $\mathbf{w}^* \in \operatorname{argmin}_{\{\mathbf{w}: \|\mathbf{w}\| \leq B\}} f(\mathbf{w})$ 
  - $f(\mathbf{w}^*)$  is a minima for  $\|\mathbf{w}\| \leq B$

## *Then:*

- If we run the GD algorithm on  $f$  for  $T$  steps with  $\eta = \sqrt{\frac{B^2}{\rho^2 T}}$ , then the output vector  $\bar{\mathbf{w}}$  satisfies:  $f(\bar{\mathbf{w}}) - f(\mathbf{w}^*) \leq \frac{B\rho}{\sqrt{T}}$

Output of GD

Minima (e.g., ERM solution)

*Demonstration  
not part of the course*

# Gradient Descent: Corollary

Theorem:

Hypothesis:  $f(\mathbf{w})$  convex  $\rho$ -Lipschitz function,  $\mathbf{w}^* \in \operatorname{argmin}_{\{\mathbf{w}: \|\mathbf{w}\| \leq B\}} f(\mathbf{w})$

Thesis :If we run the GD algorithm on  $f$  for  $T$  steps with  $\eta = \sqrt{\frac{B^2}{\rho^2 T}}$ , then the output vector  $\bar{\mathbf{w}}$  satisfies:  $f(\bar{\mathbf{w}}) - f(\mathbf{w}^*) \leq \frac{B\rho}{\sqrt{T}}$



**Corollary:**

For every  $\epsilon > 0$ , to achieve  $f(\bar{\mathbf{w}}) - f(\mathbf{w}^*) \leq \epsilon$  it suffices to run the GD algorithm for a number of iterations that satisfies  $T \geq \frac{B^2 \rho^2}{\epsilon^2}$

*Demonstration:*

- Theorem: If we run for  $T$  iterations we get that  $f(\bar{\mathbf{w}}) - f(\mathbf{w}^*) \leq \frac{B\rho}{\sqrt{T}}$
- Need  $\frac{B\rho}{\sqrt{T}} \leq \epsilon \rightarrow \sqrt{T} \geq \frac{B\rho}{\epsilon} \rightarrow T \geq \frac{B^2 \rho^2}{\epsilon^2}$

# Stochastic Gradient Descent (SGD)



SGD iterations

average of  $\mathbf{w}^{(t)}$

Example with function  
 $1.25(x + 6)^2 + (y - 8)^2$

- Computing the gradient at each step is computationally demanding  
→ in ML: need to loop through the whole training set at each iteration
- avoid using exactly the gradient  
**SGD**: take a (random) vector with expected value equal to the gradient direction

*SGD algorithm:*

$$\mathbf{w}^{(0)} \leftarrow \mathbf{0}$$

for  $t \leftarrow 0$  to  $T - 1$  do

choose  $\mathbf{v}_t$  at random from a distribution  
such that  $\mathbb{E}[\mathbf{v}_t | \mathbf{w}^{(t)}] = \nabla f(\mathbf{w}^{(t)})$

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \mathbf{v}_t$$

return  $\bar{\mathbf{w}} = \mathbf{w}^{(T)}$  (or  $\bar{\mathbf{w}} = \frac{1}{T-T_0} \sum_{t=T_0}^T \mathbf{w}^{(t)}$ )



# SGD vs GD: Notes (1)

Why should we use SGD instead of GD in machine learning applications ?

Consider the ML ERM setting:

find  $\mathbf{w}$  that minimizes  $L_S(\mathbf{w})$ , i.e.,  $f(\mathbf{w}) = L_S(\mathbf{w})$

Using **GD**:

- $\nabla f(\mathbf{w})$  depends on **all** the  $m$  pairs  $(\mathbf{x}_i, y_i) \in S$
- Need to process **all the training set at each iteration**
- Very long computation time if training set is large (as in real world ML problems)

Using **SGD**:

Need to pick  $\mathbf{v}_t$  such that  $E[\mathbf{v}_t | \mathbf{w}^{(t)}] = \nabla f(\mathbf{w}^{(t)}) = \nabla L_S(\mathbf{w}^{(t)})$

- **pick a random  $(\mathbf{x}_i, y_i) \in S \Rightarrow \mathbf{v}_t = \nabla \ell(\mathbf{w}^{(t)}, (\mathbf{x}_i, y_i))$**
- Satisfies the requirement
- Can be computed from just **a single sample** ( $\rightarrow$  much faster !!)

Same discussion apply to regularized losses and other risk minimization framework

# SGD vs GD: Notes (2)

- ❑ Much faster than GD: at each step only one sample is used for the computation
  - Specially for large training sets standard GD is slow
- ❑ Less stable trajectory
  - More "*noisy*" but could jump out of local minima
    - Advanced approaches to stabilize, e.g., momentum
  - Sometimes the final point is computed as average of a set of samples (as in the book) to account for fluctuations
    - Better to average only a set of final iterations
    - On book average of all iterations (not always smart choice)
  - Improvement to get a stable result: use an adaptive step size



# Gradient Descent: Variants

1. *Batch Gradient Descent* (standard **GD**): compute the gradient over the complete training set
2. *Mini-batch Gradient Descent*: compute the gradient over a small set of  $k$  samples
  - $k$ : parameter, mini-batch size
  - Trade-off between the two "extreme" cases **GD** and **SGD**
  - Used to train deep neural networks
3. *Stochastic Gradient Descent (SGD)*: use a single sample to estimate the gradient

# SGD: Applications in ML

Use SGD to solve ML problems :

1. Risk minimization (ERM)
2. Regularized Loss minimization (RLM)
3. Support Vector Machines (SVM)
4. Neural Networks (in NN / deep learning lectures)

# SGD for Risk Minimization (1)

Stochastic Gradient Descent (SGD) for minimizing  $L_D(\mathbf{w})$

params: Scalar  $\eta > 0$  , integer  $T > 0$

Init:  $\mathbf{w}^{(1)} = \mathbf{0}$

for  $t = 1, 2, \dots, T$

    sample  $z \sim D$

    pick  $\mathbf{v}_t = \nabla \ell(\mathbf{w}^{(t)}, z)$

    update  $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \mathbf{v}_t$

output  $\mathbf{w}^{(T)}$

- ❑ Minimize  $L_D$  directly
- ❑ Find an *unbiased* estimate of the gradient of  $L_D$
- ❑ Sample a single fresh sample and estimate the gradient with it
- ❑ Can be applied to RLM solving its target

# SGD Finds an Unbiased Estimate of the Gradient

SGD finds an unbiased estimate of the gradient of  $L_D$ :

1. Sample  $z \sim D$ :  $\mathbf{v}_t = \nabla l(\mathbf{w}^{(t)}, z)$
2.  $E[\mathbf{v}_t | \mathbf{w}^{(t)}] = E_{z \sim D}[\nabla l(\mathbf{w}^{(t)}, z)] = \nabla E_{z \sim D}[l(\mathbf{w}^{(t)}, z)] = \nabla L_D(\mathbf{w}^{(t)})$

- The gradient of the loss function  $l(\mathbf{w}^{(t)}, z)$  at  $\mathbf{w}^{(t)}$  is an unbiased estimate of the gradient of the true risk  $L_D(\mathbf{w}^{(t)})$
- Proof based on linearity of the gradient operator
- Need to sample a new fresh sample at each iteration
  - But in many optimization schemes training samples are re-used

# SGD for Risk Minimization (2)

- Consider a **convex**  $\rho$ -Lipschitz-bounded learning problem with parameters  $\rho, B$
- Then, for every  $\epsilon > 0$ , if we run the SGD method for minimizing  $L_D(\mathbf{w})$  with a number of iterations (i.e., number of examples)

$T \geq \frac{B^2 \rho^2}{\epsilon^2}$  and with  $\eta = \sqrt{\frac{B^2}{\rho^2 T}}$ , the output  $\bar{\mathbf{w}}$  of SGD satisfies:

$$\mathbb{E}[L_D(\bar{\mathbf{w}})] \leq \min_{\mathbf{w} \in \mathcal{H}} L_D(\mathbf{w}) + \epsilon$$

# SGD for $\lambda$ -strongly convex functions and RLM

- SGD for  $\lambda$ -strongly convex functions: a good strategy is to use an

adaptive step size of value  $\eta_t = \frac{1}{\lambda t}$

Idea: smaller and smaller steps while training goes on

- Details and theoretical bounds on the book, *not part of the course*

- Recall: RLM  $\rightarrow$  The associated optimization problem can be written as

$$\min_{\mathbf{w}} \left( \frac{\lambda}{2} \|\mathbf{w}\|^2 + L_s(\mathbf{w}) \right)$$

- Define  $f(\mathbf{w}) = \frac{\lambda}{2} \|\mathbf{w}\|^2 + L_s(\mathbf{w})$  : it is  $2\frac{\lambda}{2} = \lambda$ -strongly convex

- Can apply adaptive learning rate with rate  $\eta_t = \frac{1}{\lambda t}$

Not Part of the course

# SGD for RLM

- Recall:  $f(\mathbf{w}) = \frac{\lambda}{2} \|\mathbf{w}\|^2 + L_s(\mathbf{w})$ : it is  $\lambda$ -strongly convex, use  $\eta_t = \frac{1}{\lambda t}$
- Update rule can be rewritten as

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \frac{1}{\lambda t} (\lambda \mathbf{w}^{(t)} + \mathbf{v}_t) = -\frac{1}{\lambda t} \sum_{i=1}^t \mathbf{v}_i$$

*Demonstration: see next slide*

- If loss is  $\rho$ -Lipschitz, after  $T$  iterations we have that:

$$\mathbb{E}[f(\bar{\mathbf{w}})] - f(\mathbf{w}^*) \leq \frac{4\rho^2}{\lambda T} (1 + \log(T))$$

*Demonstration and details not part of the course*



# SGD for RLM (demonstration)

□ Update rule can be rewritten as

$$\begin{aligned}\mathbf{w}^{(t+1)} &= \mathbf{w}^{(t)} - \frac{1}{\lambda t} (\lambda \mathbf{w}^{(t)} + \mathbf{v}_t) \\ &= \left(1 - \frac{1}{t}\right) \mathbf{w}^{(t)} - \frac{1}{\lambda t} \mathbf{v}_t = \left(\frac{t-1}{t}\right) \mathbf{w}^{(t)} - \frac{1}{\lambda t} \mathbf{v}_t \\ &= \frac{t-1}{t} \left( \frac{t-2}{t-1} \mathbf{w}^{(t-1)} - \frac{1}{\lambda(t-1)} \mathbf{v}_{t-1} \right) - \frac{1}{\lambda t} \mathbf{v}_t \\ &= \dots\end{aligned}$$

$$= -\frac{1}{\lambda t} \sum_{i=1}^t \mathbf{v}_i$$

*We'll use for SVM*



# SGD for Soft SVM (1)

## □ Hinge loss

$$f^{hinge}(\mathbf{w}) = \max\{0, 1 - y \langle \mathbf{w}, \mathbf{x} \rangle\}$$

## □ (sub)gradient of $f^{hinge}$ at $\mathbf{w}$ :

$$\mathbf{v}^{hinge} = \begin{cases} 0 & \text{if } 1 - y \langle \mathbf{w}, \mathbf{x} \rangle \leq 0 \\ -y\mathbf{x} & \text{if } 1 - y \langle \mathbf{w}, \mathbf{x} \rangle > 0 \end{cases}$$

## □ Update Rule (for the complete soft-SVM optimization)

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \mathbf{v}^{(t)} \quad \text{or} \quad \mathbf{w}^{(t+1)} = -\frac{1}{\lambda t} \sum_{j=1}^t \mathbf{v}^{(j)}$$

- *the first equation is standard SGD*
- *the second is from the variant of SGD for  $\lambda$ -strongly convex functions*

# SGD for Soft SVM (2)

We want to solve

$$\min_{\mathbf{w}} \left( \frac{\lambda}{2} \|\mathbf{w}\|^2 + \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y \langle \mathbf{w}, \mathbf{x}_i \rangle\} \right)$$

Variant of SGD for RLM

$$\mathbf{w}^{(t+1)} = -\frac{1}{\lambda t} \sum_{j=1}^t \mathbf{v}_j$$

**Note:** it's standard to add a  $\frac{1}{2}$  in the regularization term to simplify some computations.

Algorithm:

$\theta^{(1)} \leftarrow \mathbf{0}$  ;

**for**  $t \leftarrow 1$  to  $T$  **do**

let  $\mathbf{w}^{(t)} \leftarrow \frac{1}{\lambda t} \theta^{(t)}$ ;

choose  $i$  uniformly at random from  $\{1, \dots, m\}$ ;

**if**  $y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle < 1$  **then**  $\theta^{(t+1)} \leftarrow \theta^{(t)} + y_i \mathbf{x}_i$ ;

**else**  $\theta^{(t+1)} \leftarrow \theta^{(t)}$ ;

**return**  $\bar{\mathbf{w}} = \frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)}$  ; **or**  $\bar{\mathbf{w}} = \mathbf{w}^{(T)}$

$$\theta^{(t)} = -\sum_{j=1}^{t-1} \mathbf{v}_j$$

Hinge loss: gradient is 0  
if correctly classified  
and  $-y_i \mathbf{x}_i$  if error

Not Part of the course

# SGD for Soft SVM with Kernels

We want to solve:

$$\min_{\mathbf{w}} \left( \lambda \|\mathbf{w}\|^2 + \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i \langle \mathbf{w}, \psi(\mathbf{x}_i) \rangle\} \right) (*)$$

## SGD for Solving Soft-SVM with Kernels

**Goal:** Solve Equation (\*)

**parameter:**  $T$

**Initialize:**  $\beta^{(1)} = 0$

**for**  $t = 1, \dots, T$

Let  $\alpha^{(t)} = \frac{1}{\lambda t} \beta^{(t)}$

Choose  $i$  uniformly at random from  $[m]$

For all  $j \neq i$  set  $\beta_j^{(t+1)} = \beta_j^{(t)}$

If  $(y_i \sum_{j=1}^m \alpha_j^{(t)} K(\mathbf{x}_j, \mathbf{x}_i) < 1)$

Set  $\beta_i^{(t+1)} = \beta_i^{(t)} + y_i$

Else

Set  $\beta_i^{(t+1)} = \beta_i^{(t)}$

**Output:**  $\bar{\mathbf{w}} = \sum_{j=1}^m \bar{\alpha}_j \psi(\mathbf{x}_j)$  where  $\bar{\alpha} = \frac{1}{T} \sum_{t=1}^T \alpha^{(t)}$

$$\theta^{(t)} = \sum_{j=1}^m \beta_j^{(t)} \psi(\mathbf{x}_j)$$

$$\mathbf{w}^{(t)} = \sum_{j=1}^m \alpha_j^{(t)} \psi(\mathbf{x}_j)$$

We want to solve

*Standard Soft-SVM*

$$\min_{\mathbf{w}} \left( \frac{\lambda}{2} \|\mathbf{w}\|^2 + \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y \langle \mathbf{w}, \mathbf{x}_i \rangle\} \right)$$

**Note:** it's standard to add a  $\frac{1}{2}$  in the regularization term to simplify some computations.

**Algorithm:**

$\theta^{(1)} \leftarrow \mathbf{0}$  ;

**for**  $t \leftarrow 1$  to  $T$  **do**

let  $\mathbf{w}^{(t)} \leftarrow \frac{1}{\lambda t} \theta^{(t)}$ ;

choose  $i$  uniformly at random from  $\{1, \dots, m\}$ ;

**if**  $y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle < 1$  **then**  $\theta^{(t+1)} \leftarrow \theta^{(t)} + y_i \mathbf{x}_i$ ;

**else**  $\theta^{(t+1)} \leftarrow \theta^{(t)}$ ;

**return**  $\bar{\mathbf{w}} = \frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)}$ ;

*compare with previous algorithm*

- $\mathbf{w} = \sum_{i=1}^m \alpha_i \psi(\mathbf{x}_i)$  from Representer theorem: maintain  $\alpha$  instead of  $\mathbf{w}$
- **Theorem:** the new procedure and the old one applied on the feature space (i.e., replacing  $\mathbf{x}$  with  $\psi(\mathbf{x})$ ) lead to the same results (*no demonstration*)


# SGD: Issues

- ❑ The selection of the learning rate  $\eta$  is a critical point
  - If  $\eta$  too small: the optimization is **stable** but the convergence can be **very slow**
  - If  $\eta$  too large: the convergence is **fast** but the optimization can be **very unstable**
- ❑ Simple solution: use adaptive learning rates, e.g.,
  - Progressively reducing the learning rate according to a pre-defined schedule
  - *Example* : for RLM optimization with SGD  $\eta_t = \frac{1}{\lambda t}$  is used
  - However these approaches requires rules and thresholds to be defined in advance and thus are difficult to adapt to different problems
- ❑ Additionally, the same learning rate applies to all parameter updates
  - The various parameters have different behaviors and the learning rate could be too fast for some and too slow for others
  - Sometimes better not to update all parameters to the same extent, but perform a larger update for some and smaller for others

# Momentum



SGD without momentum



SGD with momentum

- SGD has troubles (i.e., it oscillates) in areas where the surface curves much more steeply in one dimension than in another (which are common around local optima)
- **Momentum**: the update is the linear combination of previous gradient and new one
  - The momentum parameter  $\gamma$  is usually set to 0.9 or a similar value

$$\begin{aligned} \text{○ } \mathbf{v}^{(t)} &= \gamma \mathbf{v}^{(t-1)} + (1 - \gamma) \nabla L(\mathbf{w}^{(t)}) \\ \text{○ } \mathbf{w}^{(t+1)} &= \mathbf{w}^{(t)} - \eta \mathbf{v}^{(t)} \end{aligned}$$

- It helps accelerate SGD in the relevant direction and dampens oscillations

Using momentum is like pushing a ball down a hill. The ball accumulates momentum as it rolls downhill, becoming faster and faster on the way. The same thing happens to our parameter updates: the momentum term increases for dimensions whose gradients point in the same directions and reduces updates for dimensions whose gradients change directions. As a result, we gain faster convergence and reduced oscillation

# Advanced SGD schemes

- ❑ **Adagrad** adapts the learning rate for each parameter independently
  - It performs smaller updates (i.e. low learning rates) for parameters associated with frequently occurring features
  - It performs larger updates (i.e. high learning rates) for parameters associated with infrequent features
- ❑ **Adadelta** (improved version of ADAgrad)
- ❑ **RMSprop** (improved version of ADAgrad)
- ❑ **Adam** (Adaptive Moment Estimation)
  - It also computes adaptive learning rates for each parameter
  - It combines ideas from Adagrad and momentum
  - Whereas momentum can be seen as a ball running down a slope, Adam behaves like a heavy ball with friction, which thus prefers flat minima in the error surface