

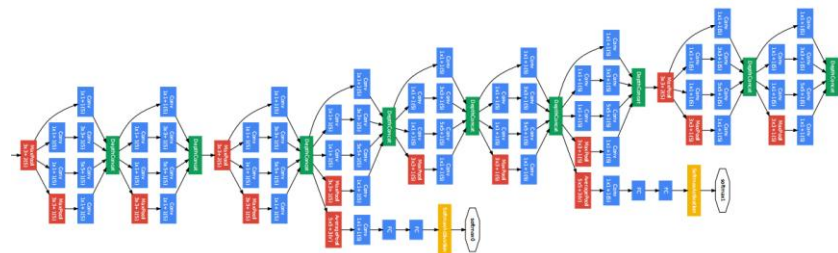
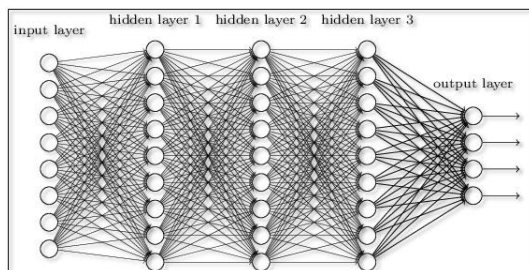
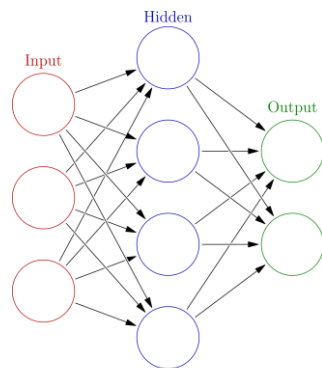
Convolutional Neural Networks

Machine Learning 2023-24

Slides P. Zanuttigh

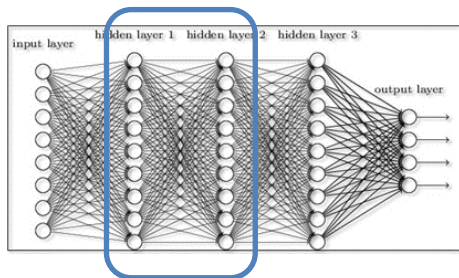
Some slides Ming Li

Recall: Artificial Neural Networks



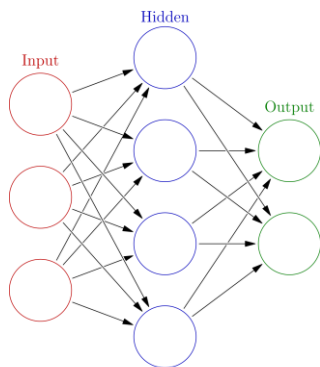
- Model of computation inspired by the structure of neural networks in the brain
- Large number of basic computing devices (**neurons**) connected to each other
- Represented with directed graphs where the nodes are the neurons and the edges corresponds to the links between the neurons
- Proposed in 1940-50
- First practical applications in the 80-90 but practical results were lower than SVM and other techniques
- **From 2010 on deep architectures with impressive performances**

Recall: Feedforward NN



Feedforward network: the graph has no edges

It is typically organized into layers: **each neuron** takes in input the output of **all neurons** from the previous layer

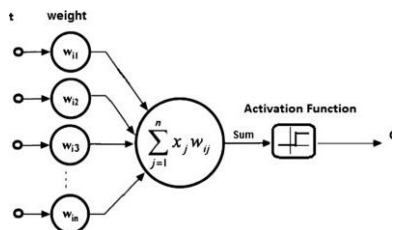


Notation: NN: $G=(V,E)$

- V : neurons $|V|$: size of the network
- E : connection between neurons (directed edges)
- $w: E \rightarrow \mathbb{R}$ weight function over the edges

Each neuron:

1. Takes in input the sum of the outputs of the connected neurons weighted by the edge weights
2. Applies to it a simple scalar function (activation function, σ)



NN Training Algorithm

BackPropagation algorithm with SGD

Input: training data $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$

Output: NN weights $w_{ij}^{(t)}$

Initialize $w_{ij}^{(t)}, \forall i, j, t$;

for $s \leftarrow 0, 1, 2, \dots$ **do**

// until convergence

pick (\mathbf{x}_k, y_k) at random from training data; // SGD

compute $v_{t,j}, \forall j, t$;

// forward propagation

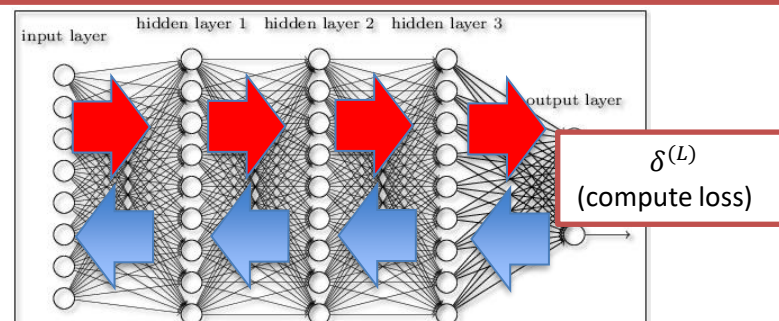
compute $\delta_j^{(t)}, \forall j, t$;

// backward propagation

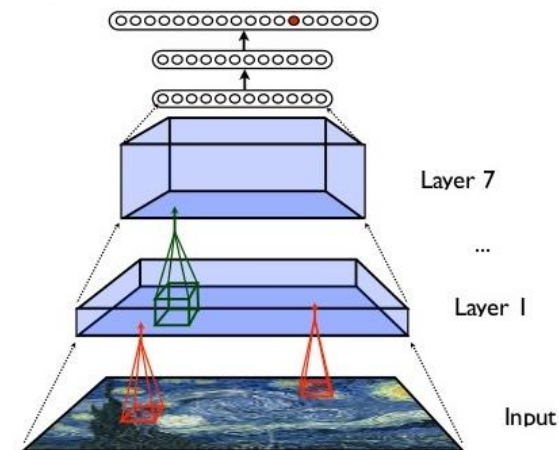
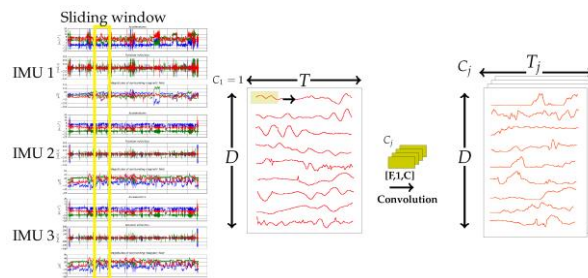
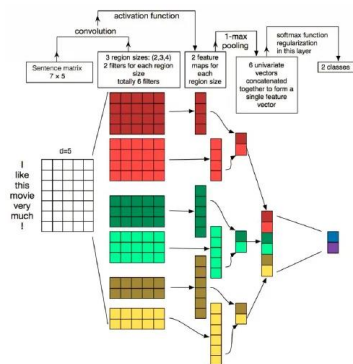
$w_{ij}^{(t)[s+1]} = w_{ij}^{(t)[s]} - \eta v_{t-1,i} \delta_j^{(t)} \forall i, j, t$;

// update weights

if converged then return $w_{ij}^{(t)}, \forall i, j, t$;



Issues of Fully Connected Feedforward Networks



Two main issues in the NN model we have seen up to now:

1. **Each** neuron of layer $t-1$ connected with **each** neuron of layer t
 → *huge number of edges/weights (quadratic w.r.t. the number of neurons in each layer)*
2. *The domain structure is not taken into account*
 - The model does not consider that a neuron can be "closer" (→ more related) to some neurons and less to others
 - **Some domains have a structure**
 - E.g., grid of pixels in an image, sequence of samples in an audio signal, letters of a word in a text, ...
 - Need to capture the fact that a pixel in an image is more related to the close pixels than to the far apart ones or a letter in a text is more related to letters of the same word than to the ones 10 pages ahead !
 - Interesting features are often local, shift-invariant and deformation-invariant
 - By simply placing data in a vector → loose spatial or temporal structure

→ *Need to update the NN model*

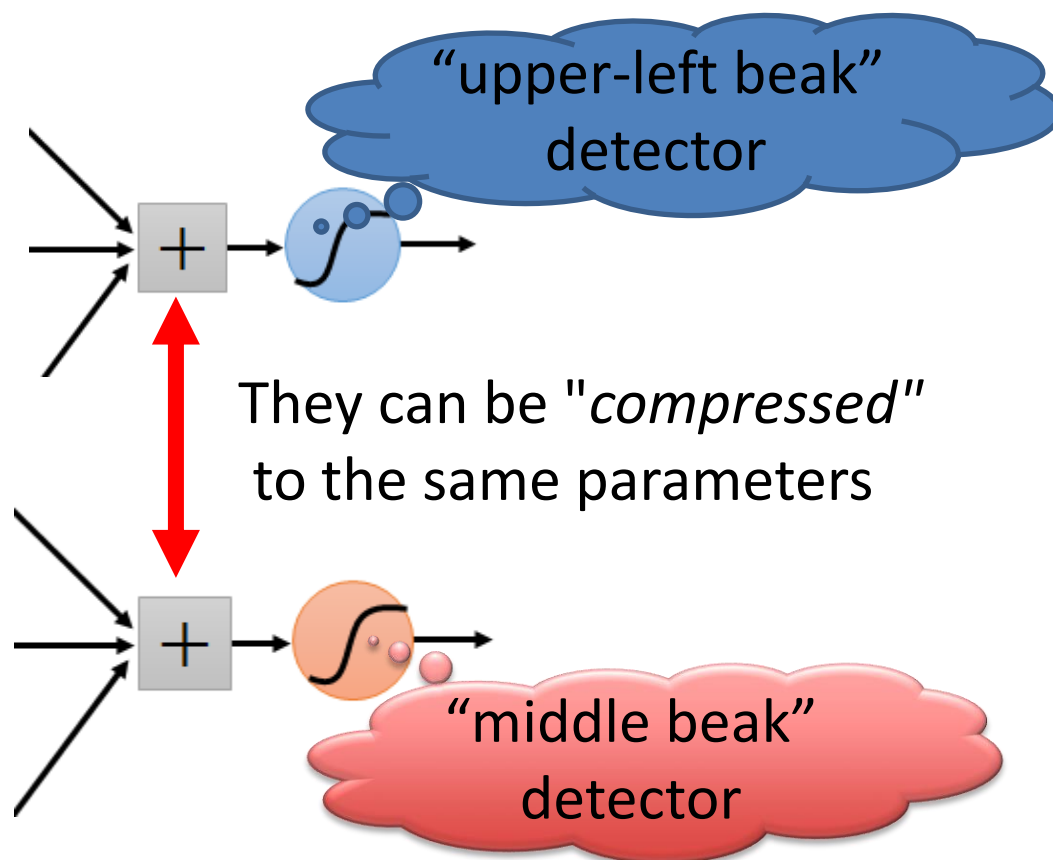
Example: Learning an Image (1)

Recognize patterns that are much smaller than
the whole image

Can represent a small region with fewer parameters

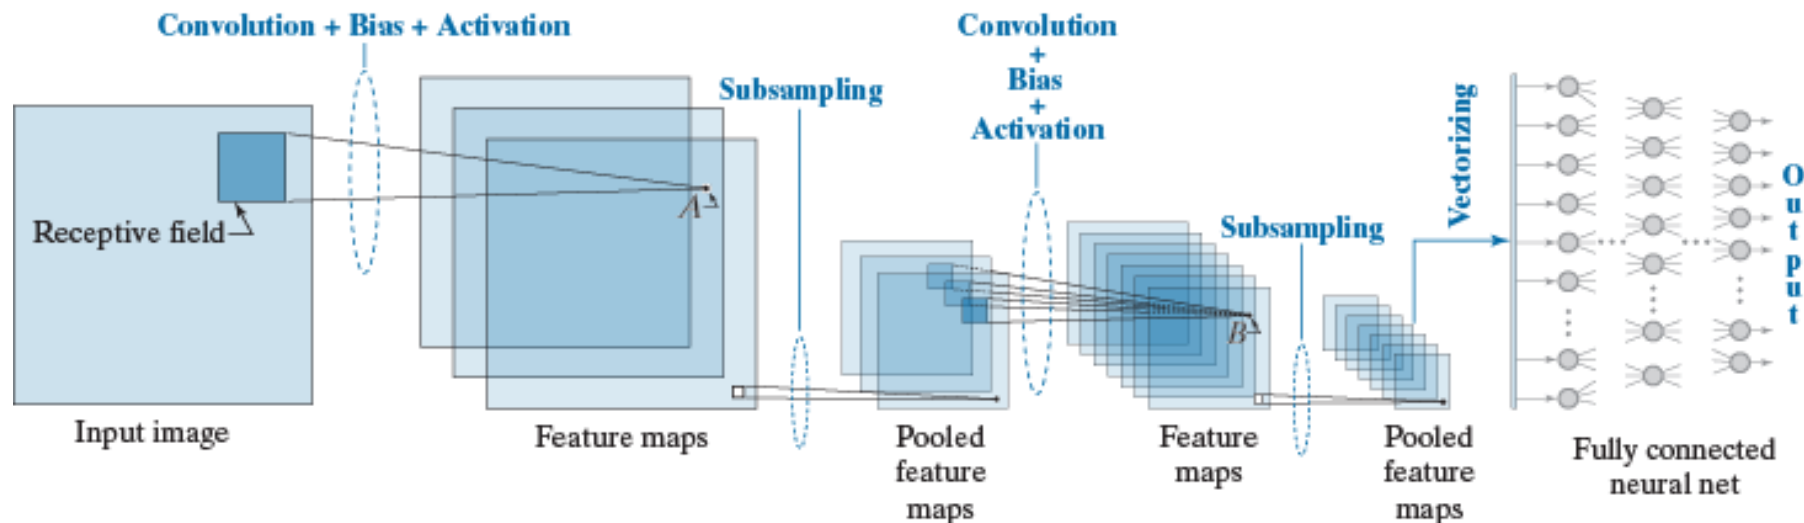


Example: Learning an Image (2)



- The same pattern can appear in different places
- Similar detectors in different regions share similar parameters
- What about training some “small” detectors and let each detector “move around” ?

From NNs to Convolutional Neural Networks

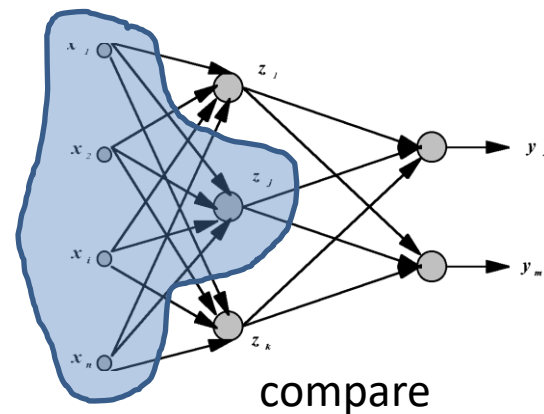
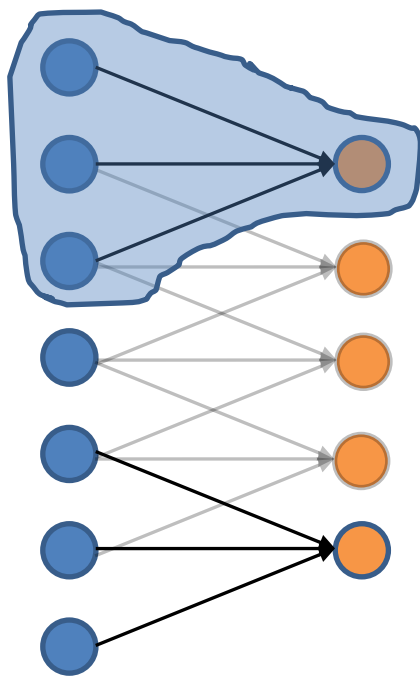


Convolutional Neural Network (CNN)

1. *Local connectivity*: receptive field for each neuron
2. *Shared ("tied") weights*: spatially invariant response
3. *Multiple feature maps*
4. Subsampling (*pooling*)

Convolutional NNs

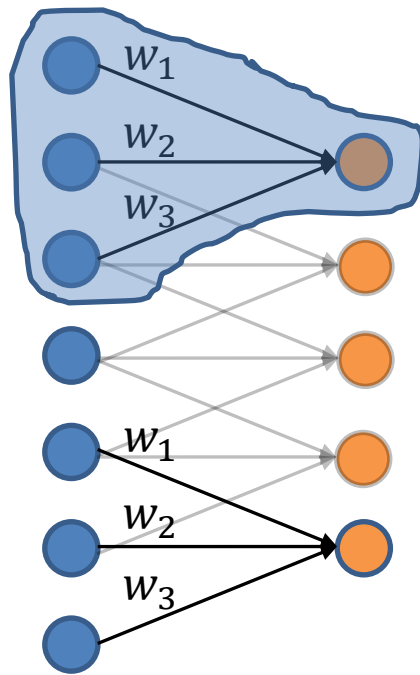
1. Local connectivity



- Each orange unit (neuron) is **only** connected to **neighboring** blue units

Convolutional NNs

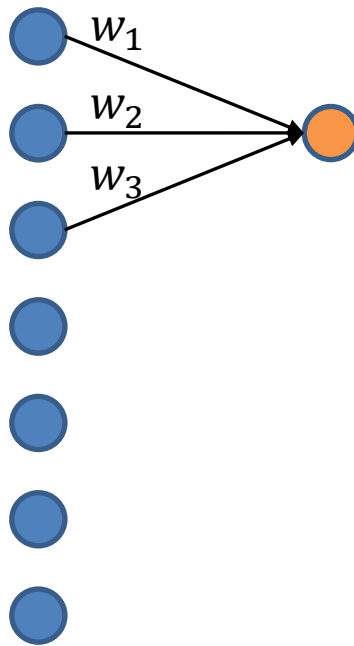
2. Shared (“tied”) weights



- All orange units **share** the same parameters w
- Each orange unit computes the **same function**, but with a **different input window**

Convolutional NNs

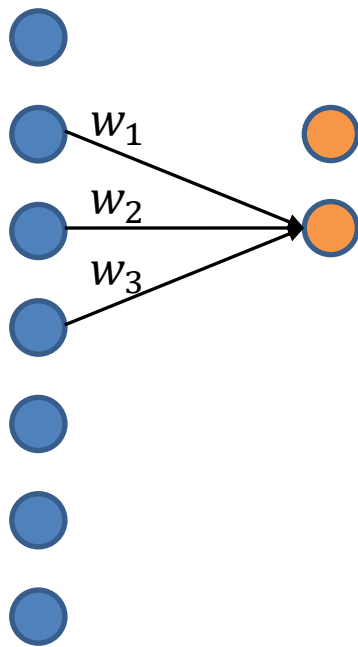
□ Convolution with 1-D filter: $[w_1, w_2, w_3]$



- All orange units **share** the same parameters \mathbf{w}
- Each orange unit computes the **same function**, but with a **different input window**

Convolutional NNs

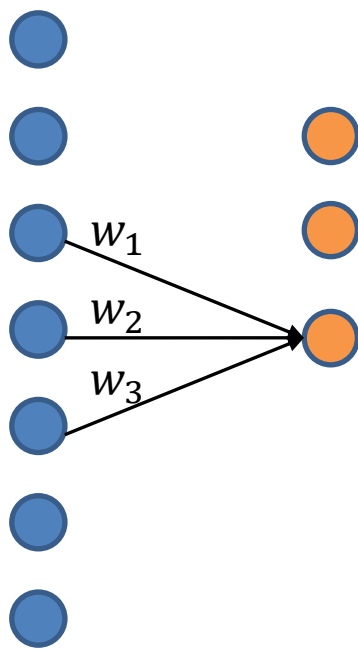
□ Convolution with 1-D filter: $[w_1, w_2, w_3]$



- All orange units **share** the same parameters \mathbf{w}
- Each orange unit computes the **same function**, but with a **different input window**

Convolutional NNs

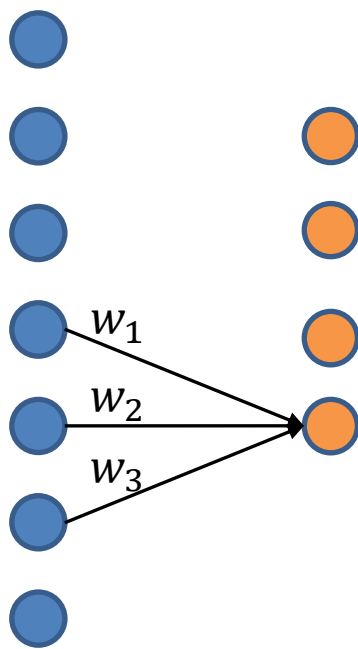
□ Convolution with 1-D filter: $[w_1, w_2, w_3]$



- All orange units **share** the same parameters \mathbf{w}
- Each orange unit computes the **same function**, but with a **different input window**

Convolutional NNs

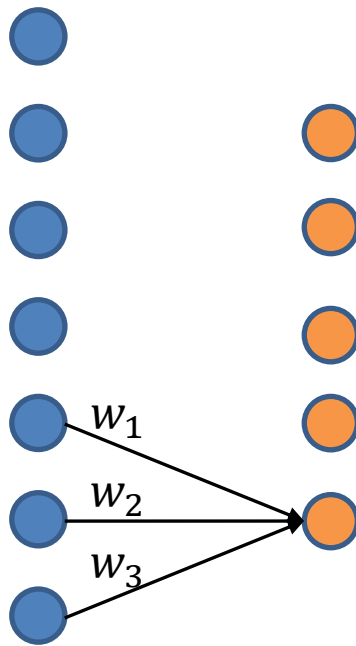
□ Convolution with 1-D filter: $[w_1, w_2, w_3]$



- All orange units **share** the same parameters \mathbf{w}
- Each orange unit computes the **same function**, but with a **different input window**

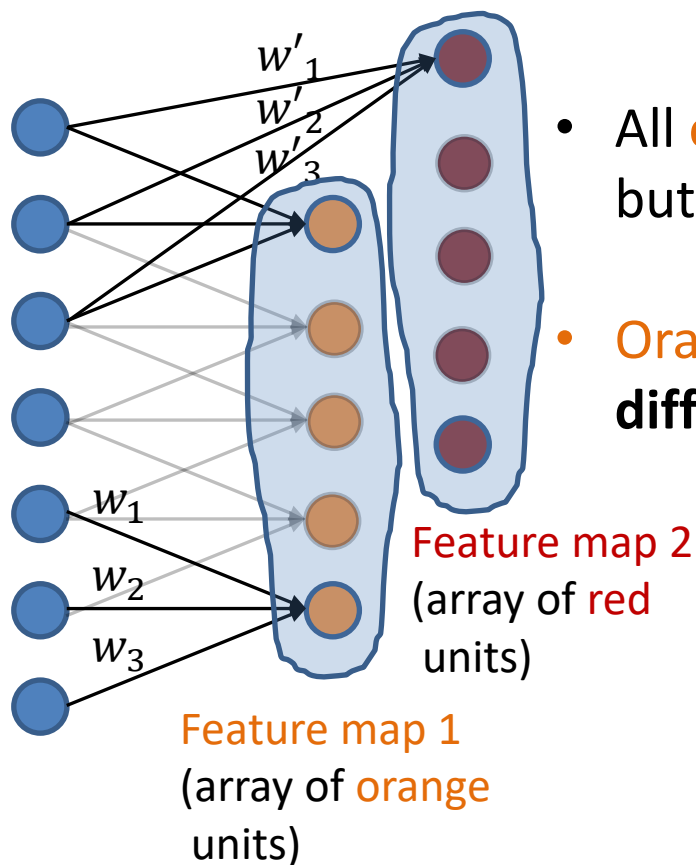
Convolutional NNs

□ Convolution with 1-D filter: $[w_1, w_2, w_3]$



- All orange units **share** the same parameters \mathbf{w}
- Each orange unit computes the **same function**, but with a **different input window**

3. Multiple feature maps



- All **orange** units compute the **same function** but with a **different input windows**
- **Orange** and **red** units **compute different functions**

Convolution

?	?	?				
?	1	0	0	0	0	1
?	0	1	0	0	1	0
	0	0	1	1	0	0
	1	0	0	0	1	0
	0	1	0	0	1	0
	0	0	1	0	1	0

6 x 6 matrix

Convolution at boundaries:

- Stop before → reduced output size
- Use padding to extend input size

These are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1
(feature map 1)

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2
(feature map 2)

⋮

Each filter detects a small pattern (3 x 3)

Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 matrix

Dot
product



3

-1

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Convolution

If **stride=2**

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 matrix

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3

-3

Stride: allows the filter to move in steps of multiple samples (alternative to pooling to reduce resolution)

Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 matrix

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Convolution

stride=1

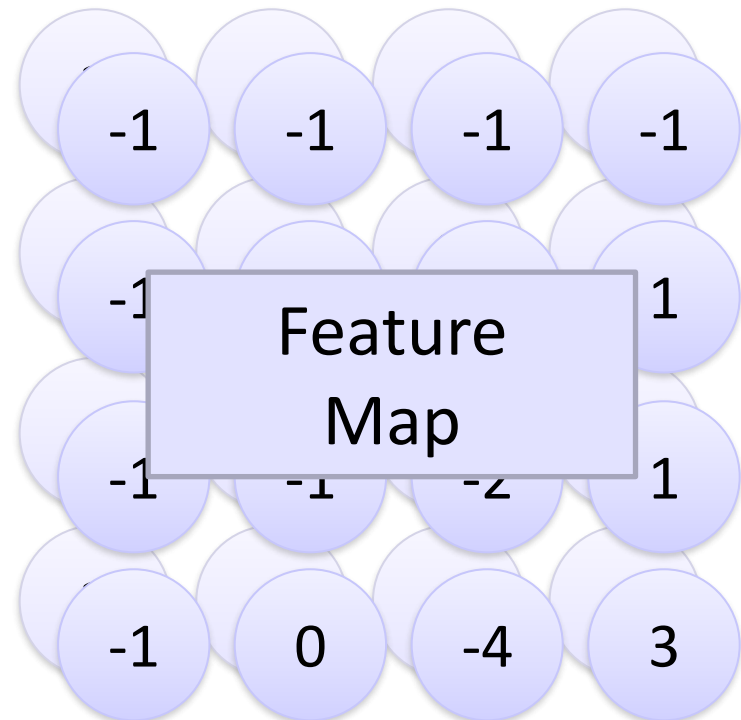
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 matrix

-1	1	-1
-1	1	-1
-1	1	-1

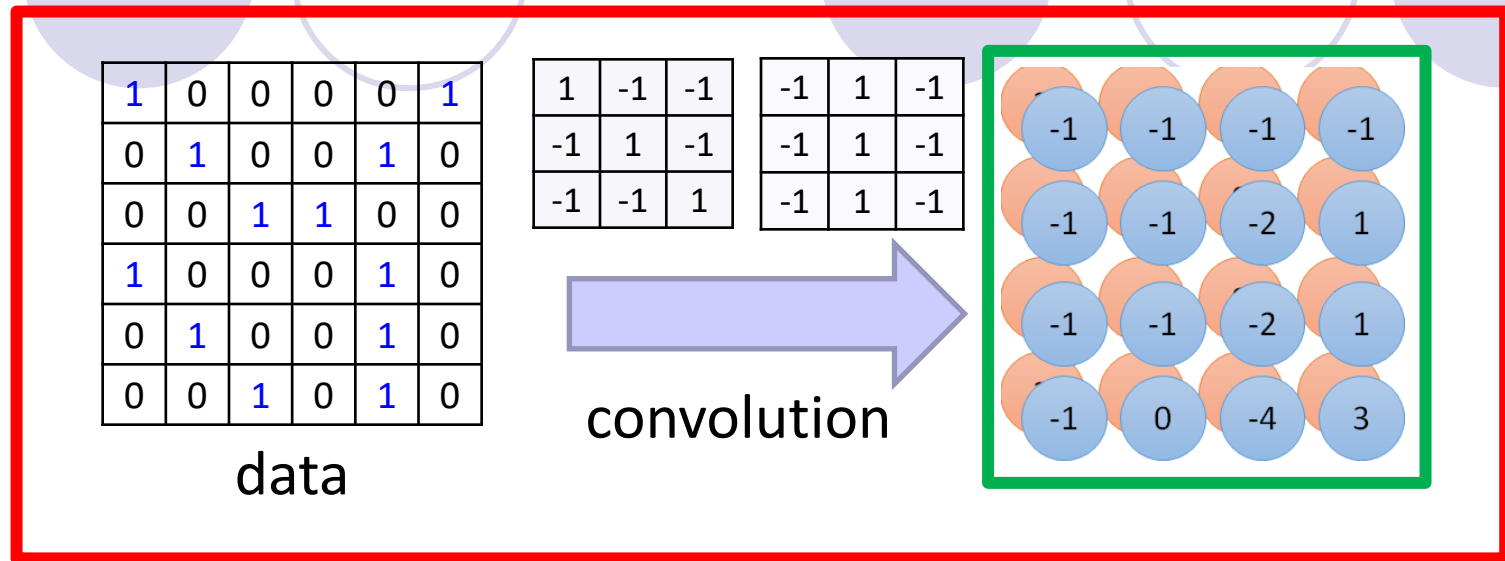
Filter 2

Repeat this for each filter



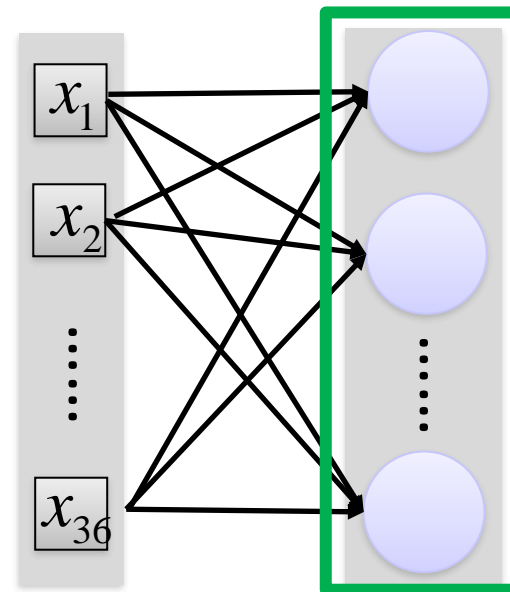
Two 4 x 4 images
Forming 2 x 4 x 4 matrix

Convolution v.s. Fully Connected

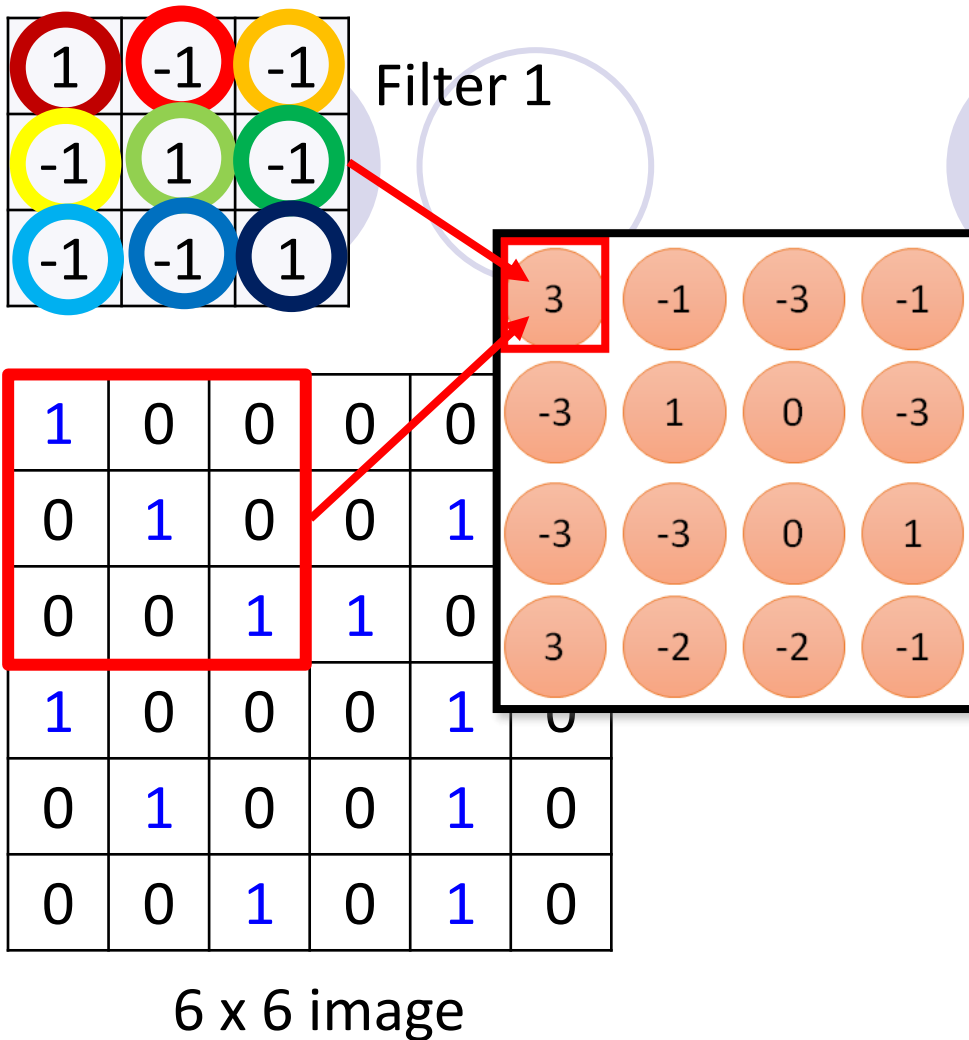


Fully-
connected

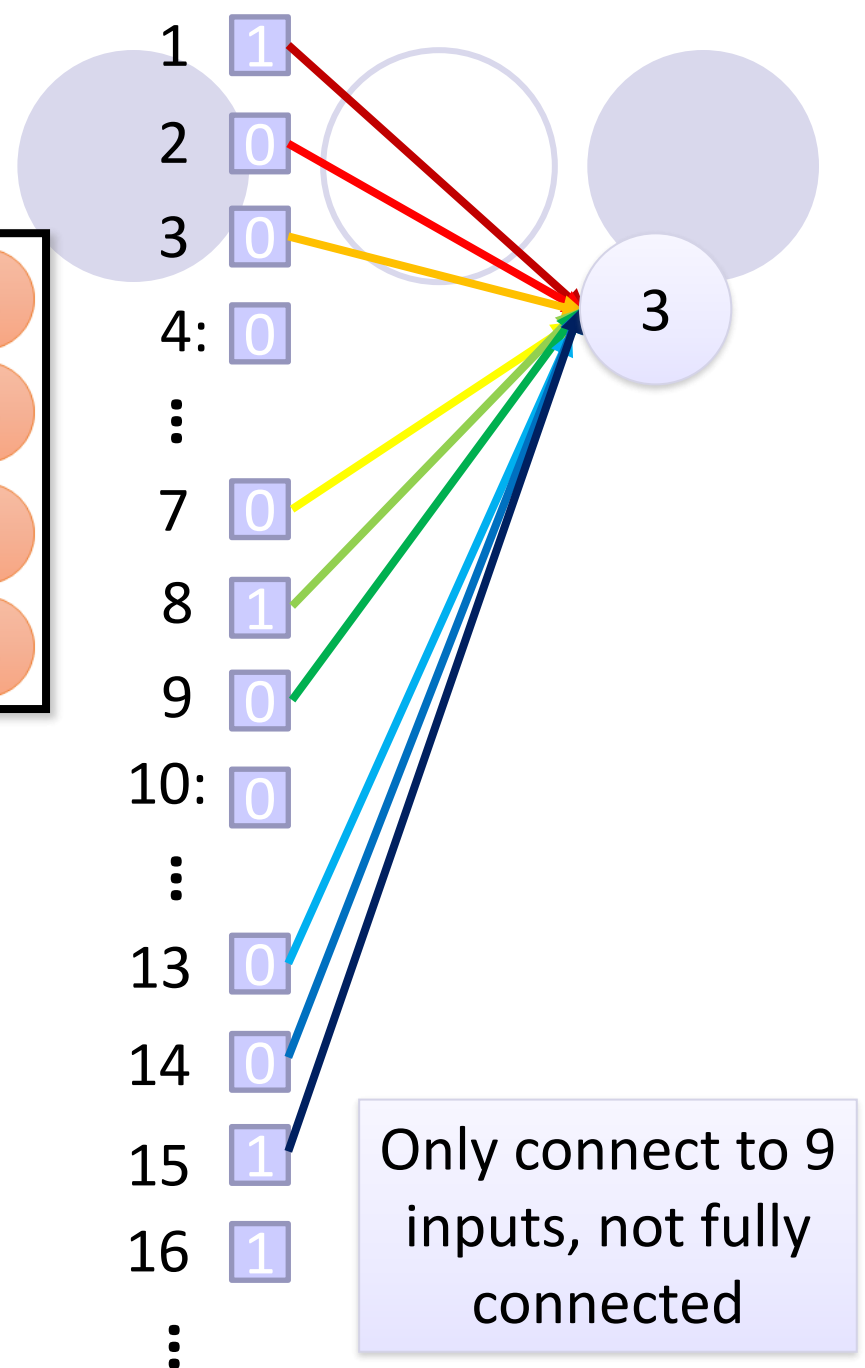
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

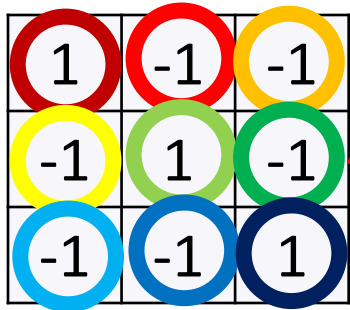


Notice: in convolutional NN #params depends only on convoution size, in fully-connected NN it depends on the data matrix size

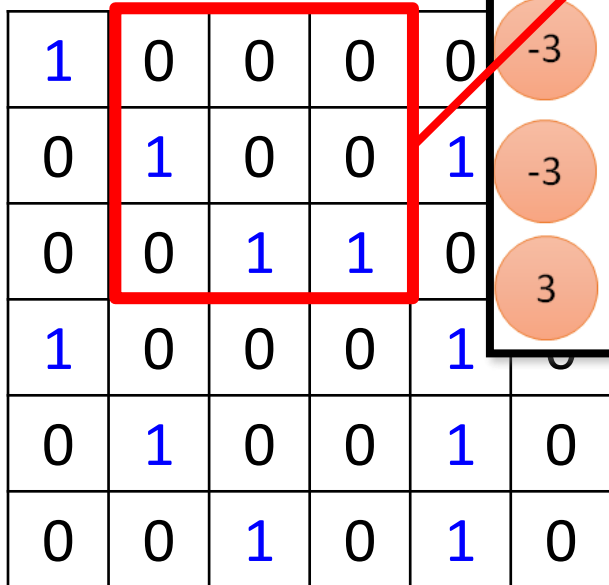


Convolutional model:
fewer parameters!





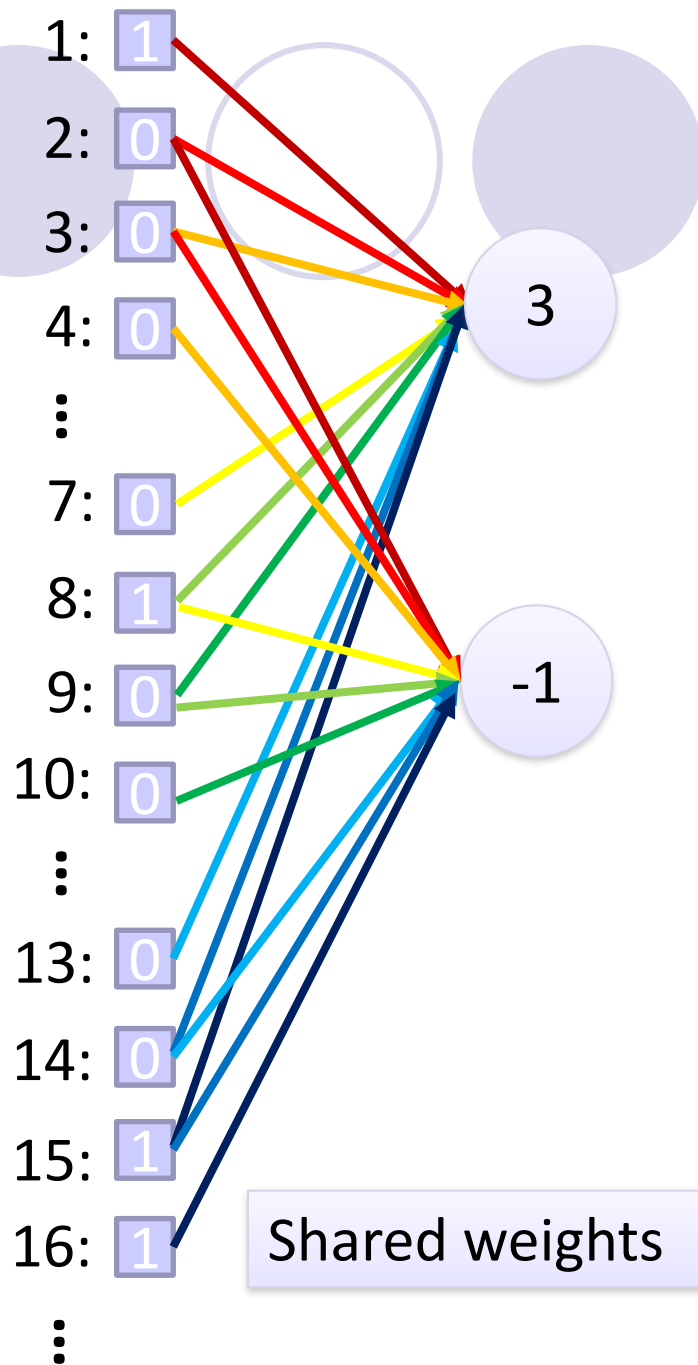
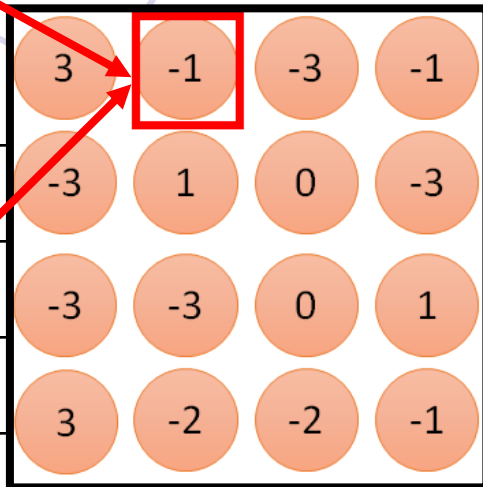
Filter 1



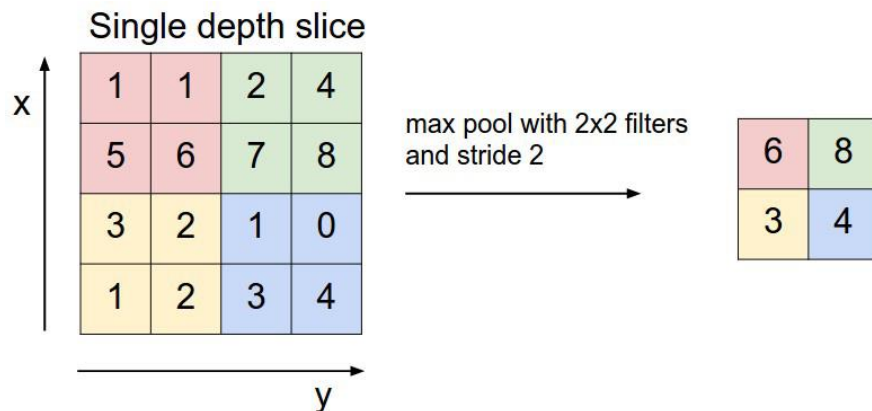
6 x 6 image

Convolutional model:
Fewer parameters

Shared weights:
Even fewer parameters!



Pooling Layer



- ❑ Reduce resolution → next convolutional layer is applied at a larger scale
- ❑ Originally introduced to reduce the computational burden and the memory requirements...
- ❑ ...but turned out to be crucial to improve performance in many applications since it increases the receptive field of the inner layers
- ❑ Adds some deformation invariance too
- ❑ *Max Pooling* is the most common example of such layer: it works very well, it is quick, and can be efficiently implemented in hardware

Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3

Idea of max pooling: preserve the strongest responses

Why Pooling ?

Image example:

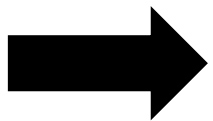
Subsampling pixels will not change the object

bird



Subsampling

bird



We can subsample the pixels to make images smaller and use fewer parameters to characterize them

However, this is not the only reason for using pooling...

Pooling and Receptive Size

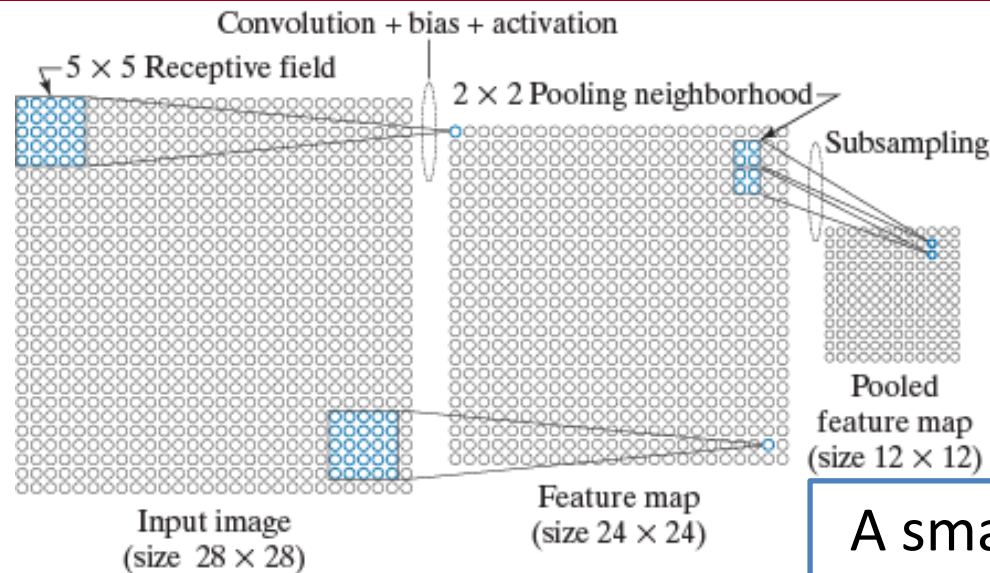


FIGURE 12.41

Top row: How the sizes of receptive fields and pooling neighborhoods affect the sizes of feature maps and pooled feature maps.

A small convolution window after pooling corresponds to a larger area in previous layers

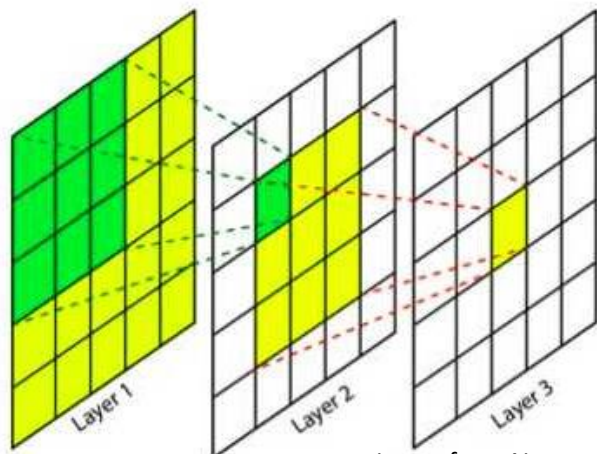
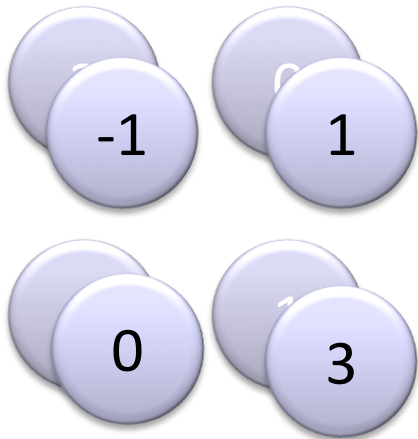


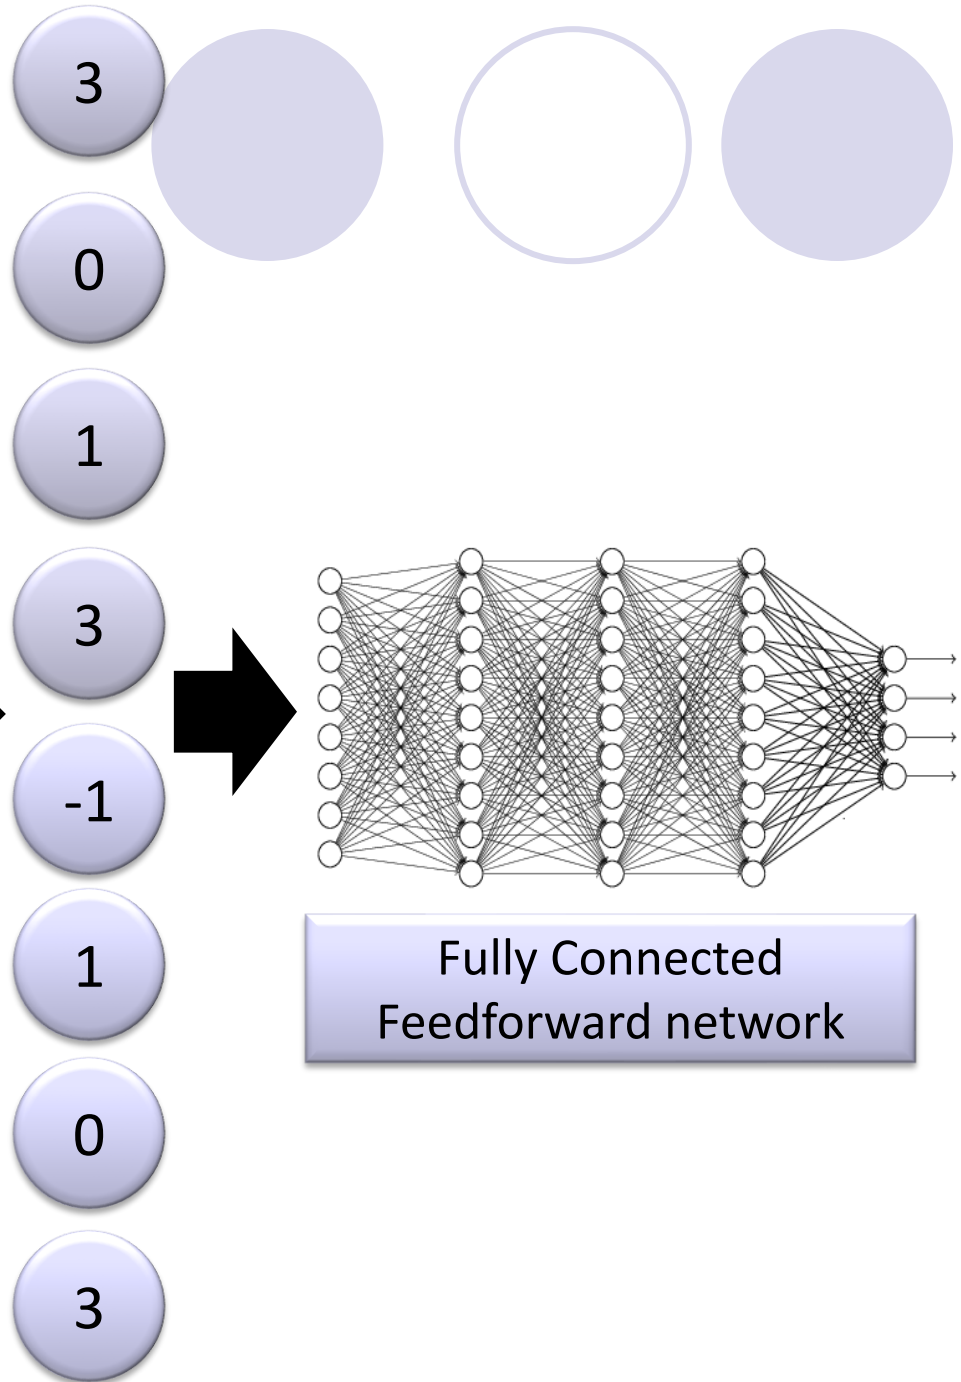
Image from AI summer



Flattening and Fully Connected Layer at the End



Flattened



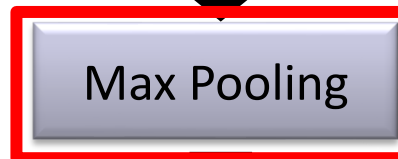
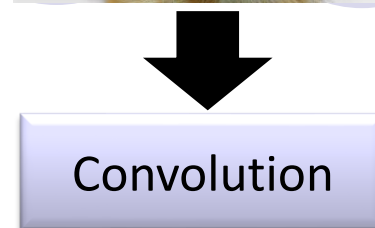
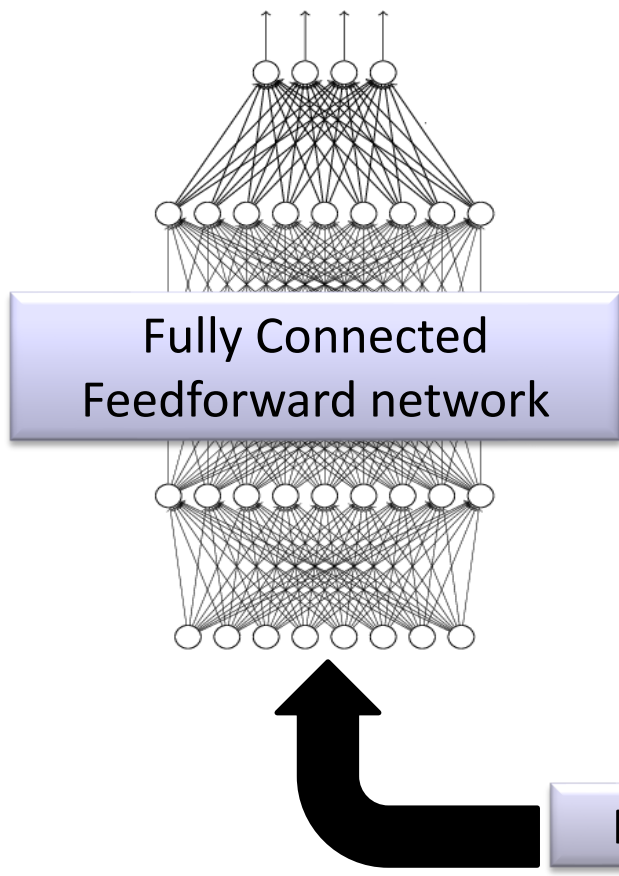
Fully Connected
Feedforward network

Final stage in classification task:
flatten (i.e., reshape to 1D) the data
and apply a set of fully connected
layers to get the final output

Baseline CNN model



cat dog

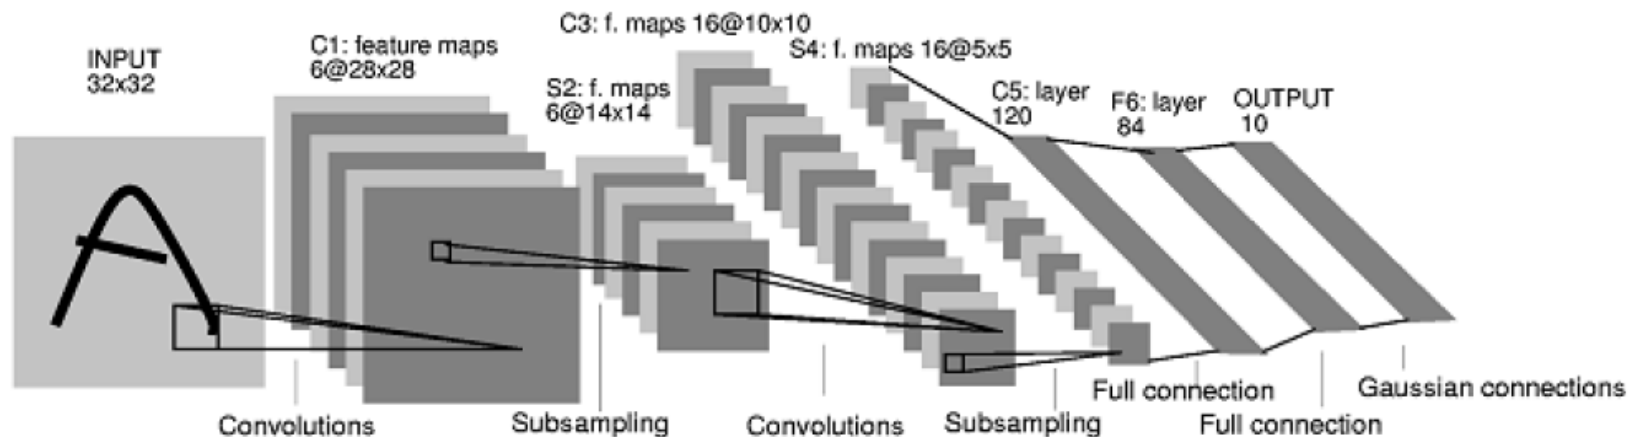


Can repeat many times



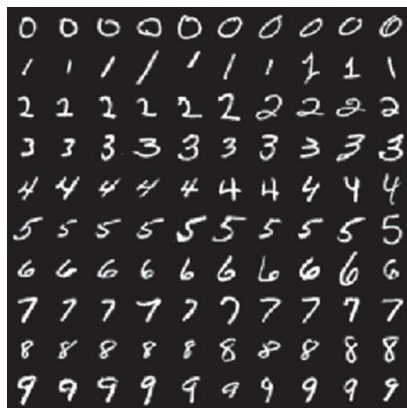
After several pooling stages the data size is much smaller→the fully connected model is affordable

Convolutional Networks



- ❑ **Hierarchical representation:** low level features in the first layer, then moving to higher and higher abstraction levels
- ❑ **Weight sharing:** huge reduction of complexity w.r.t. a fully connected network
- ❑ **CNN model "compresses" a fully connected network** in various ways:
 - Reducing the number of connections
 - Shared weights on the edges
 - Max pooling further reduces the complexity

Example: a Simple CNN

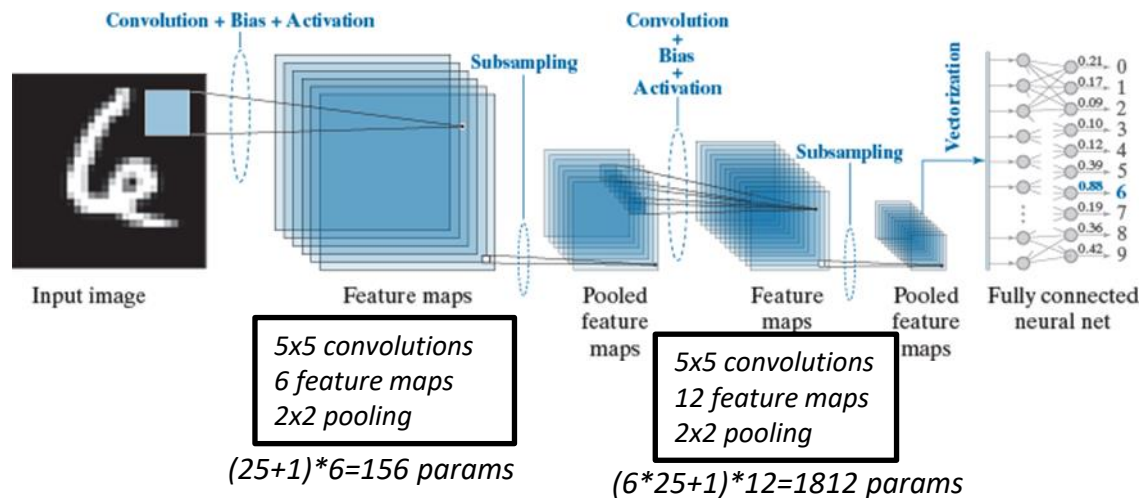


MNIST dataset

70k 28x28 images

Handwritten digits

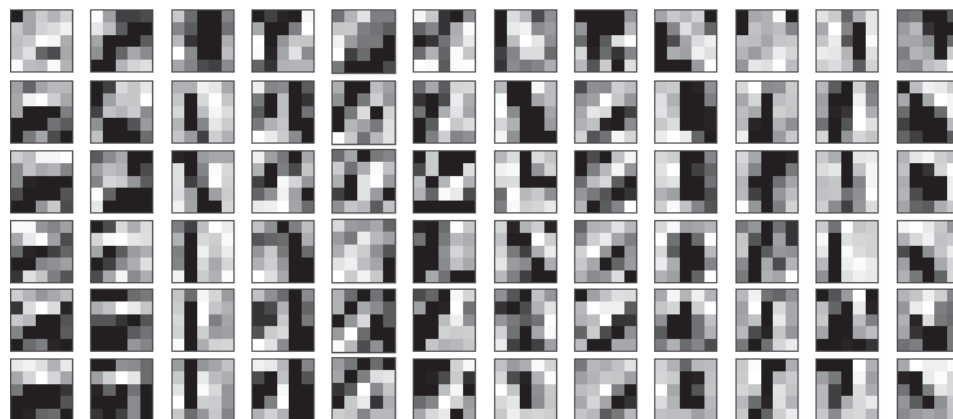
Task: classify into 10 classes



Layer 1



Layer 2



Example: Feature Maps

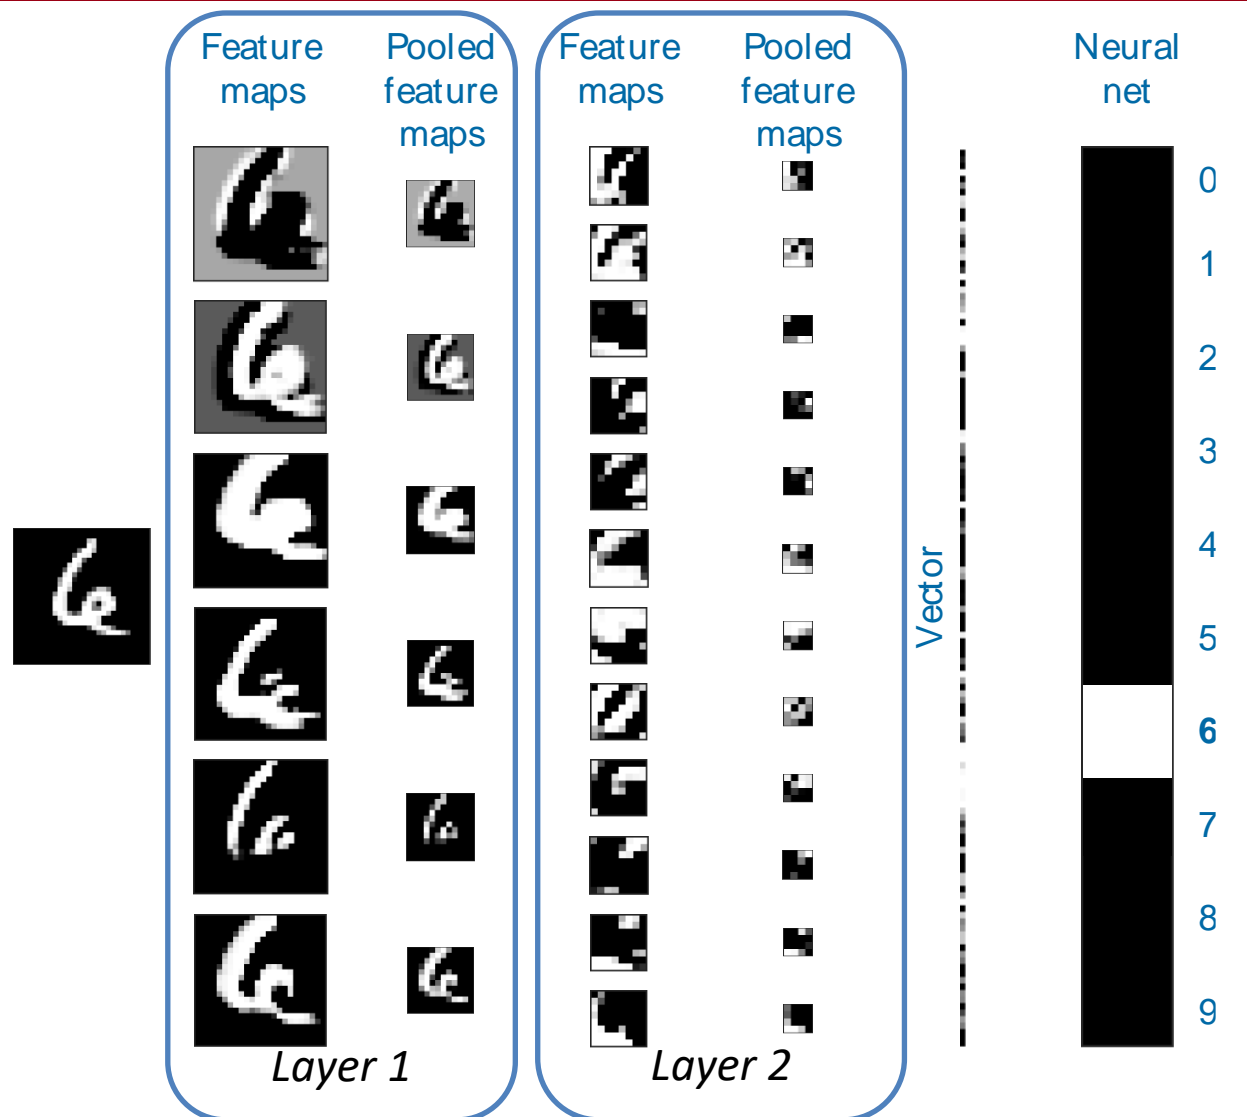
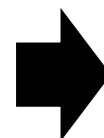


FIGURE 12.44
Visual summary
of an input image
propagating
through the CNN
in Fig. 12.42. Shown
as images are all the
results of
convolution
(feature maps) and
pooling (pooled
feature maps) for
both layers of the
network. (Example
12.17 contains more
details about this
figure.)

Example: AlphaGo



Neural
Network

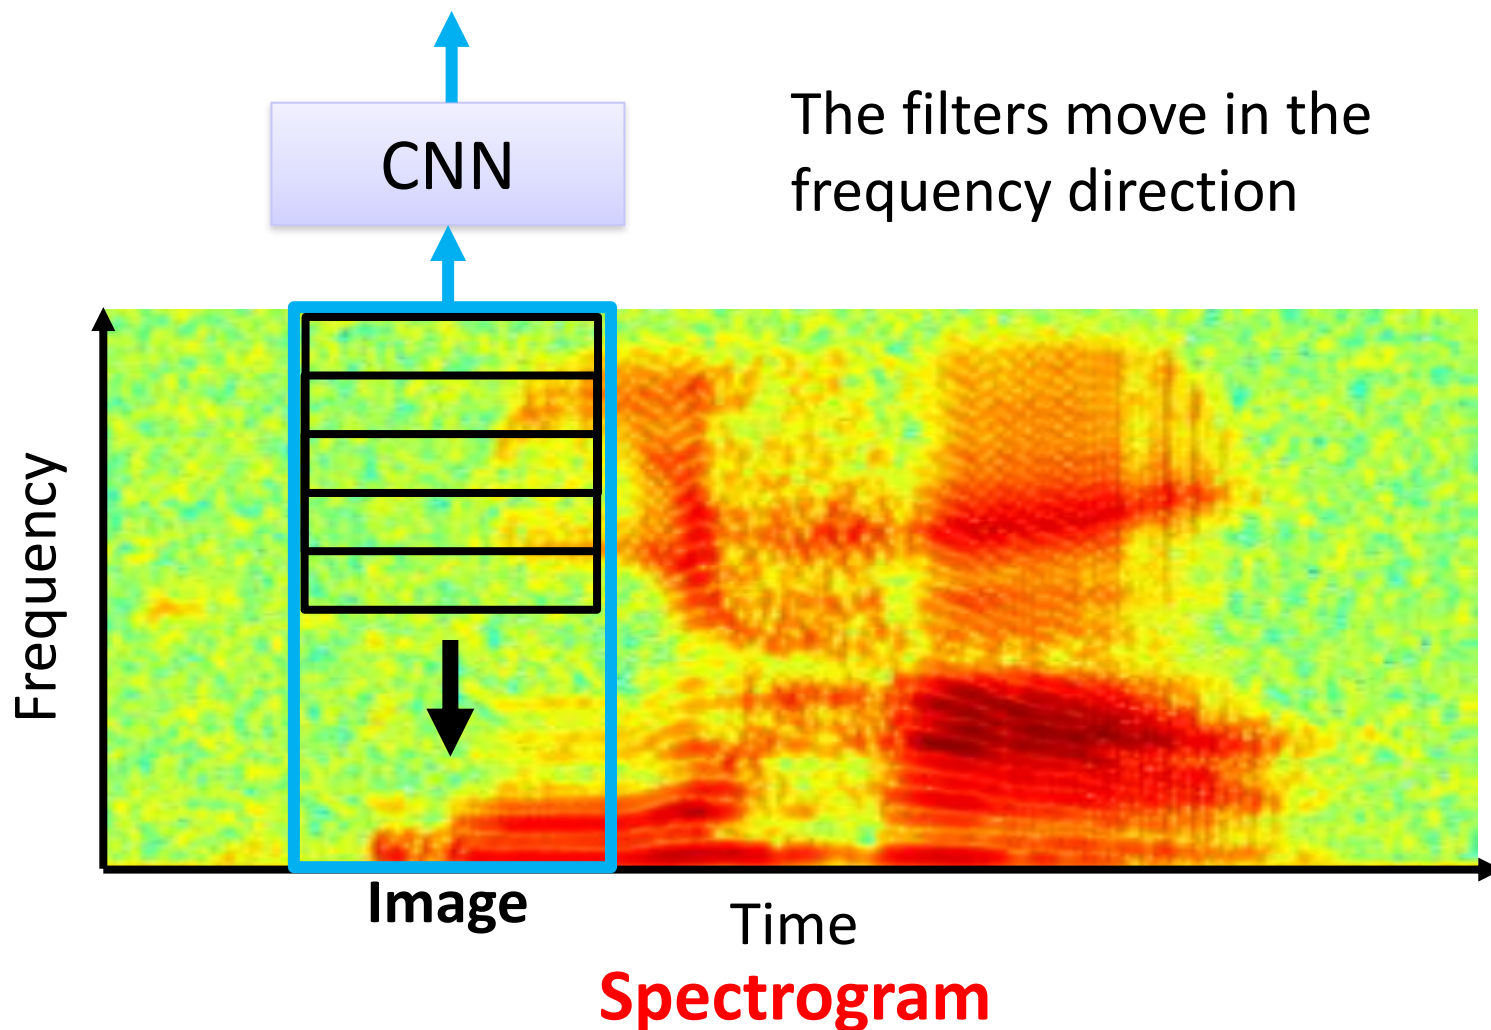


Next move
(19 x 19 positions)

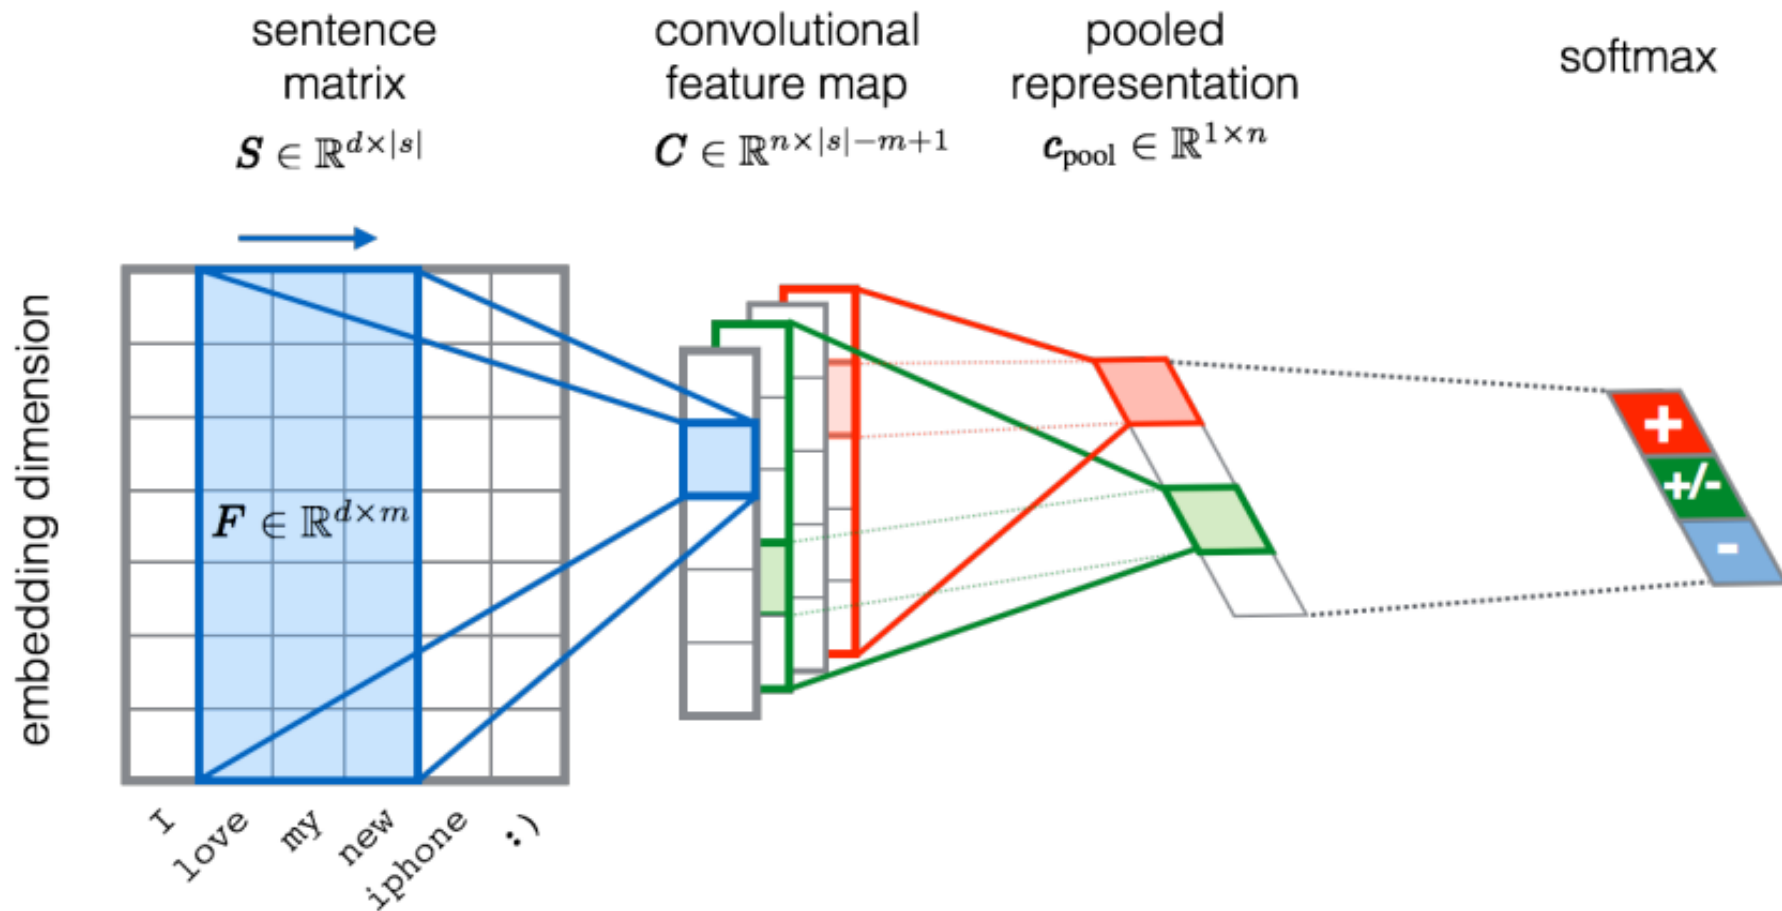
Fully-connected feedforward network
can be used

But CNN performs much better

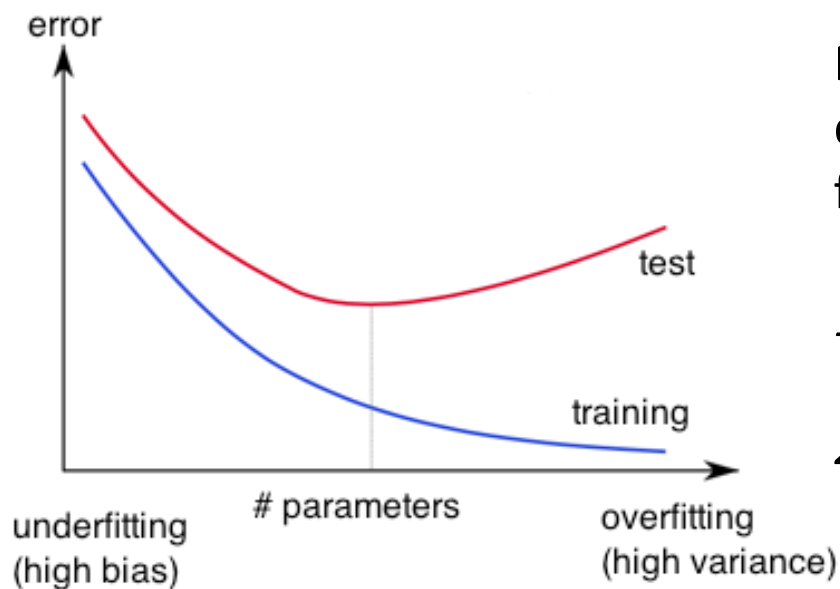
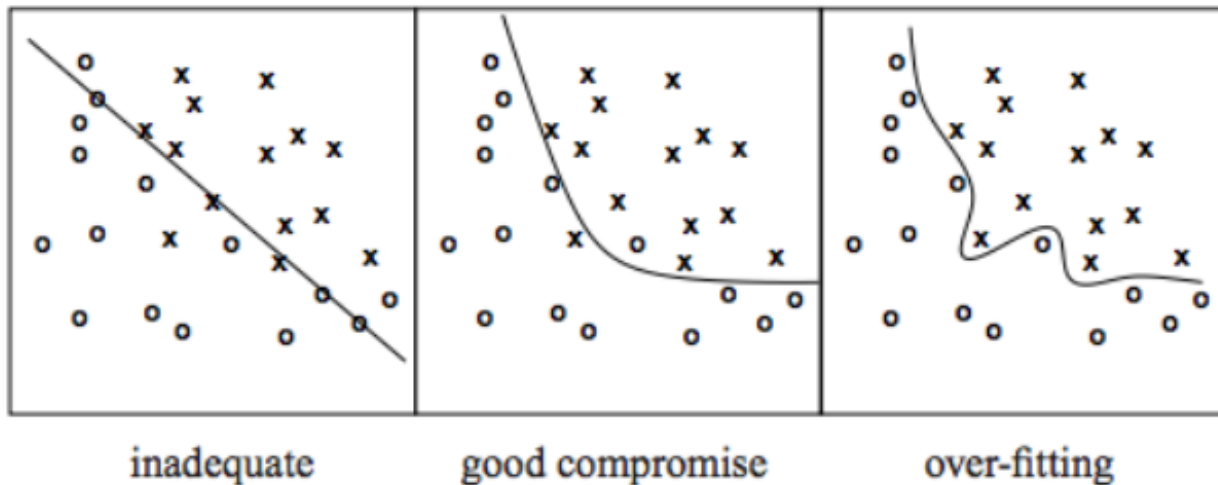
Example: CNNs in Speech Recognition



Example: CNNs in Text Classification



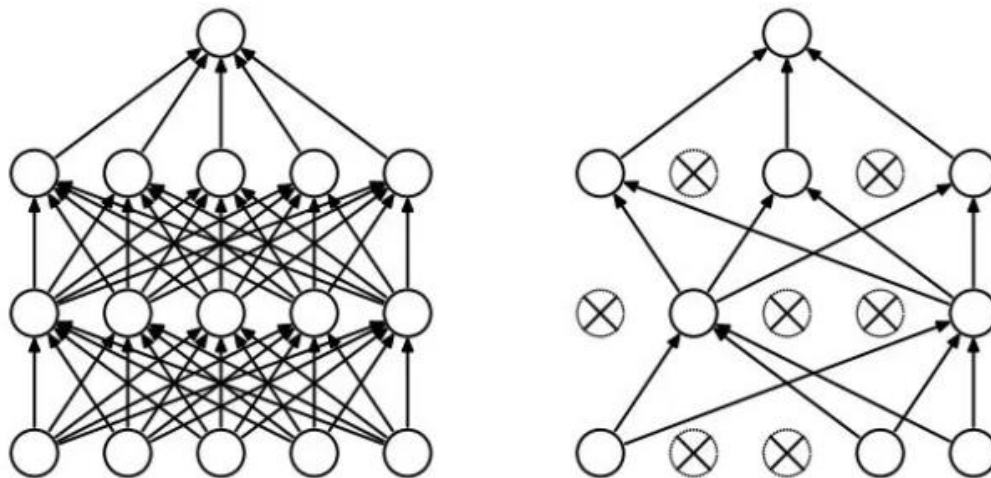
A Critical Issue: Overfitting



NN: Learned hypothesis may fit the training data very well, even outliers (noise), but could fail to generalize to new examples (test data)

1. Do not use a **too complex network** if training data is limited
2. Various techniques can be used to deal with this problem

Avoid Overfitting: Dropout



Dropout

- Randomly drop neurons (along with their connections) during training
- Each unit retained with fixed probability p , independent of other units
- *Hyper-parameter* p to be chosen (tuned)
- At each step the network is trained with only a subset of the neurons
- Avoid that the output depends "too much" on a single neuron
- Typically applied only to some layers (e.g., fully connected at the end)
- More stable / less risk of overfitting

Avoid Overfitting: Regularization

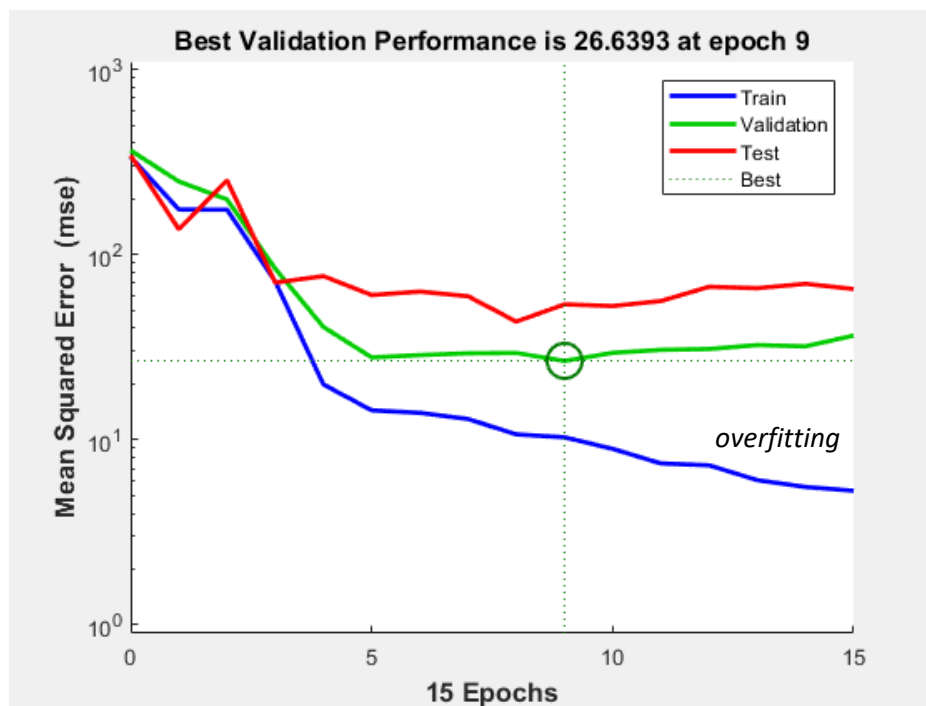
$$\text{L1 Regularization: } J_{reg}(\mathbf{w}) = J(\mathbf{w}) + \frac{\lambda}{m} \sum_{i,j,t} |\mathbf{w}_{ij}^{(t)}|$$

$$\text{L2 Regularization: } J_{reg}(\mathbf{w}) = J(\mathbf{w}) + \frac{\lambda}{m} \sum_{i,j,t} (\mathbf{w}_{ij}^{(t)})^2$$

Regularization

- Regularization term added to the loss function
- Penalizes big weights and reduces risk of overfitting
- Regularization parameter λ determines how relevant regularization is during gradient computation
- Big $\lambda \rightarrow$ big penalty for big weights (higher training error but less overfitting)
- **L1** or **L2** regularization can be used

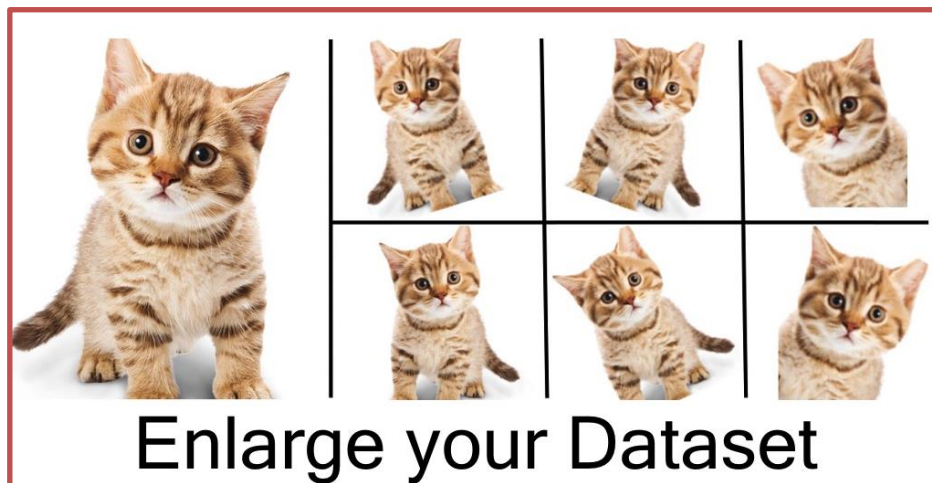
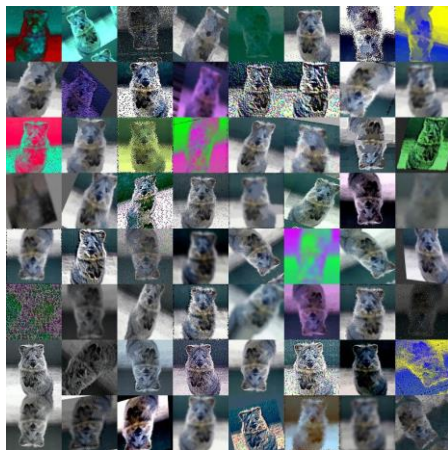
Avoid Overfitting: Early Stopping



Early-stopping

- Use validation error to decide when to stop training
- Stop when monitored loss has not improved after n subsequent epochs
- Parameter " n " is called patience

Avoid Overfitting: Data Augmentation



Data Augmentation

- ❑ Add a little bit of variance to the data to "virtually" increase number of training samples
 - Notice that new samples are correlated with the original ones, not the same as having more samples
- ❑ Artificially add noise
- ❑ Apply random transformations (depend on data type)
 - Crop part of the data
 - Resize/rescale data
 - Rotate
 - Add noise
 - Custom transformations depending on data type
 - e.g., for images: flip horizontally, adjust hue, contrast and saturation
 - e.g., for audio: time stretching, pitch shifting, amplitude scaling, reverberation

Gradient Descent for NN (1)

- ❑ Huge number of parameters, very challenging optimization
- ❑ A variety of optimization algorithms have been proposed (**recall GD lecture**)

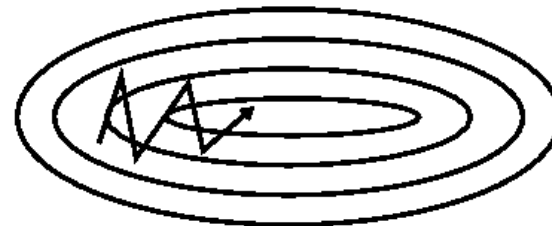
1. *Basic Gradient Descent (GD)*

- Computes the gradient of the cost function w.r.t. to the parameters for the entire training dataset
- Need to calculate the gradients for the whole dataset to perform just one update
- Can be **very slow** and is intractable for datasets that don't fit in memory

2. *Stochastic Gradient Descent (SGD)*

- Performs a parameter update for each training example
- It is usually **much faster** but performs frequent updates with a high variance and can be **unstable**
- SGD's fluctuation, on the one hand, enables it to jump to new and potentially better local minima
- On the other hand, this ultimately complicates convergence to the exact minimum, as SGD will keep overshooting

Gradient Descent for NN (2)



3. *Mini-batch gradient descent:*

- ☐ Compromise between GD and SGD: performs an update for every mini-batch of n training examples
- ☐ Reduces the variance of the parameter updates, leading to **more stable** convergence
- ☐ Can make use of highly optimized matrix computations in state-of-the-art deep learning libraries
- ☐ Common mini-batch sizes range between a few items and 256, but can vary for different applications

4. *Momentum:*

- ☐ Helps accelerate SGD in the relevant direction and dampens oscillations
- ☐ It does this by adding a fraction of the update vector of the past step to the current update vector

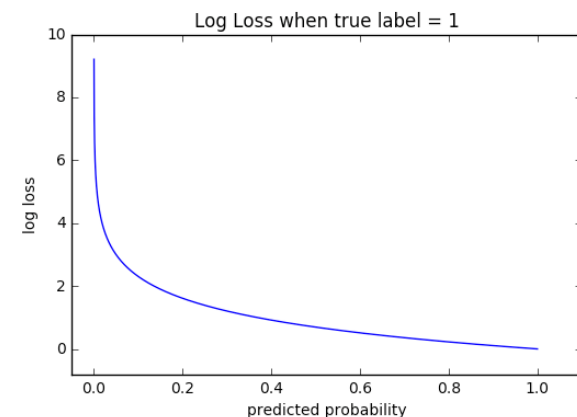
5. *Adam (Adaptive Moment Estimation)*

- ☐ Commonly used method that computes adaptive learning rates for each parameter

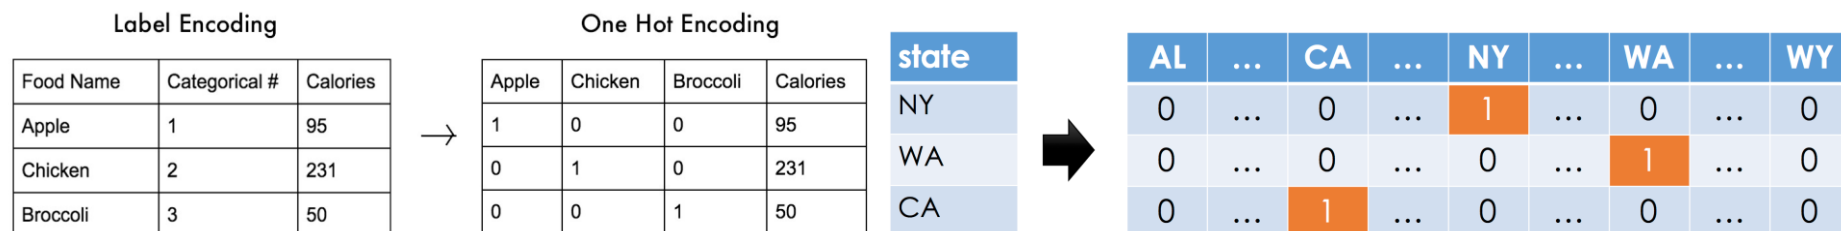
... and many others !!!!

Loss Function: Cross Entropy

- ❑ For classification tasks the *cross entropy* is commonly used in place of the 0-1 loss
- ❑ For binary classification: $L(f(\mathbf{x}), y) = -y \log(f(\mathbf{x})) - (1 - y) \log(1 - f(\mathbf{x}))$
- ❑ The optimal $f(\mathbf{x})$ minimizing this loss function is $f(\mathbf{x}) = P(y = 1 \mid \mathbf{x})$
 - We are training the neural net output to *estimate conditional probabilities*
- ❑ Note that the expression works if $f(\mathbf{x})$ is *strictly* between 0 and 1
 - An undefined or infinite value would otherwise arise
 - To achieve this, the *sigmoid* is commonly used as activation for the output layer
- ❑ The function is convex
 - Gradient descent (e.g., SGD) works better



Extension to Multi-Class



❑ One-hot encoding

- Output: vector \mathbf{y} with one component for each class
- $y_i = 1$ if sample in class i , $y_i = 0$ otherwise
- Avoid having some classes "closer" to others as when using class index
- Increases output data dimensionality

❑ Extension of cross-entropy to multi-class

- Labels one-hot encoded, vector function \mathbf{f} to be estimated
- $f_i(\mathbf{x})$ = estimated probability that \mathbf{x} belong to class i

$$L(\mathbf{f}(\mathbf{x}), \mathbf{y}) = - \sum_i y_i \log(f_i(\mathbf{x}))$$

In Practice: Many DL Tools.....

- ❑ Many deep learning frameworks
- ❑ Supported by large research entities and companies
- ❑ Optimized for GPU computing



Tensorflow (Google)



Keras: higher level framework for easier implementation



Caffe (University of Berkley)



PyTorch (Meta)



Microsoft Cognitive Toolkit

... and many others